

CSCI317 Database Performance Tuning

Denormalizations

Dr Janusz R. Getta

School of Computing and Information Technology -
University of Wollongong

Denormalizations

Outline

Simplifications

Migration of identifiers

Migration of attributes from “one” to “many” side

Simplifications

Simplification of a conceptual schema is performed in the following four steps:

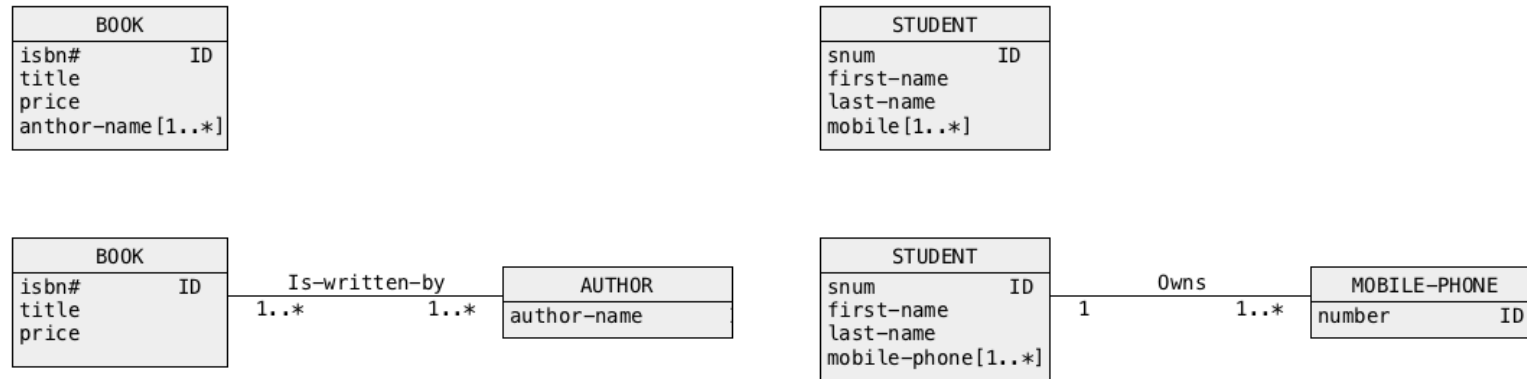
- Elimination of multivalued attributes
- Elimination of association classes
- Elimination of link attributes
- Elimination of many-to-many associations

An outcome of **simplification** is a conceptual schema that contains only **one-to-many** and **one-to-one associations**, **qualifications**, and **generalizations**

A **simplified conceptual schema** is an intermediate form of design between a **conceptual schema** and the **relational schemas**

Elimination of multivalued attributes

Elimination of multivalued attributes is performed in the same way as in a process of logical design




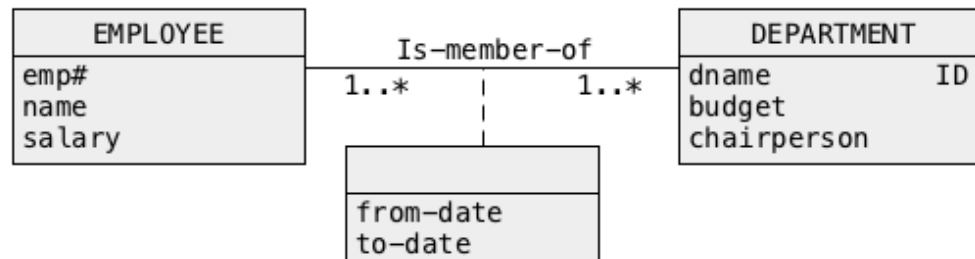
The values of multivalued attribute are promoted to objects

A multivalued attribute is replaced with a class of objects and either one-to-many or many-to-many association depending on the semantics of a multivalued attribute

A new class is described by a single valued attribute with the same name as multivalued attribute


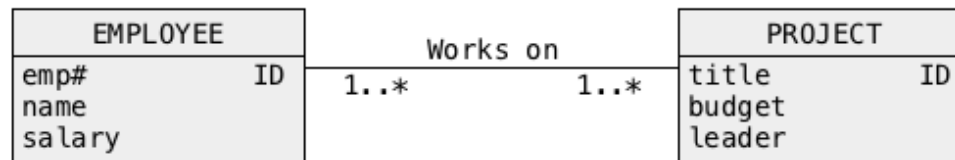
Elimination of link attributes

Elimination of **link attributes** is performed by promotion of **link attributes** to **association classes** (to be eliminated in the next step)



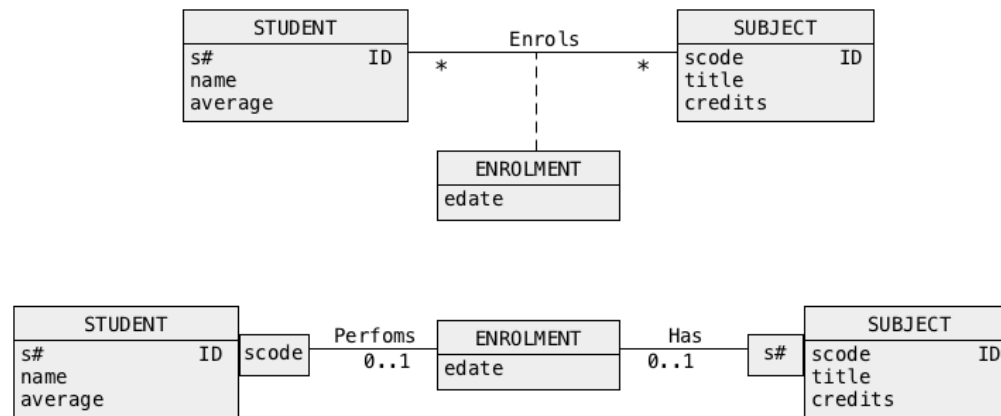
Elimination of many-to-many associations

Many-to-many association is transformed into the same **many-to-many** association and empty **association class** that can be eliminated in the next step



Elimination of association classes

Elimination of **association classes** is performed by the replacement of **many-to-many association** and **association class** attached with **one-to-many** and **many-to-one** associations



An **association class** becomes a **class of objects** involved in **one-to-many** and **many-to-one** associations and qualifications

In a general case **many-to-many association** is equivalent to **one-to-many** and **many to many-to-one** associations with a **class of objects** representing **many-to-many** association

Denormalizations

Outline

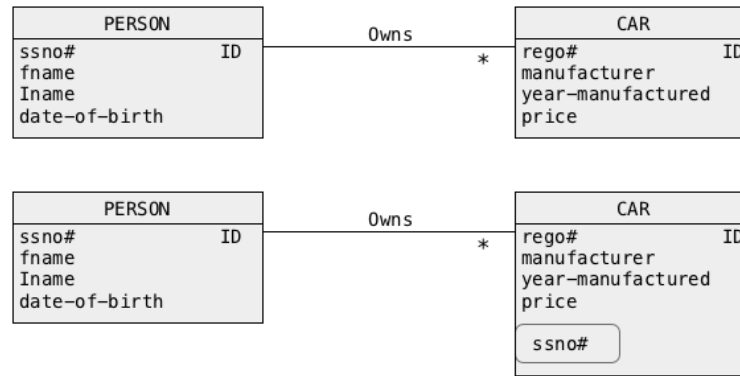
Simplifications

Migration of identifiers

Migration of attributes from “one” to “many” side

Migration of identifiers

To **implement an association** we have to copy an **identifier** from "**one**" side to "**many**" side of the **association** (and NOT the opposite !)



Performance related observations:

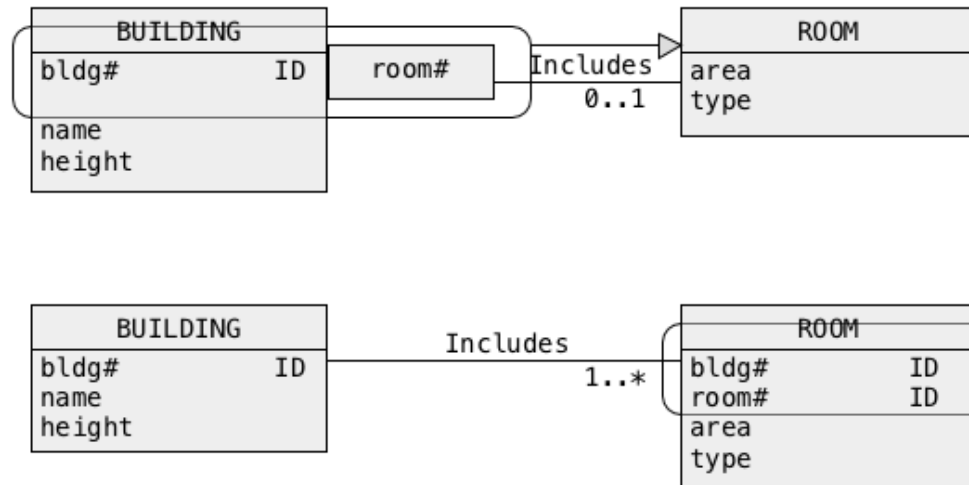
A long composite identifier allows for longer relational schemas without losing normalization

With a long composite foreign key probability that query can be computed without join is higher

Artificial identifiers (surrogate keys) contribute to join operations

Migration of identifiers

If an **identifier** of a class on "**many**" side of an association is determined through a **qualification** then both **identifier** from "**one**" side and **qualification attributes** are copied from "**one**" side to "**many**" side of the **association** (and NOT the opposite !)



Migration of identifiers

Advantages

- Migration of identifiers allows for "link-by-value" implementation of associations

Disadvantages

- Migration of identifiers creates redundancies inevitable for "link-by-value" implementation of associations

Comments

- "Link-by-reference" is the other implementation of association (object-oriented, object-relational database models)

Denormalizations

Outline

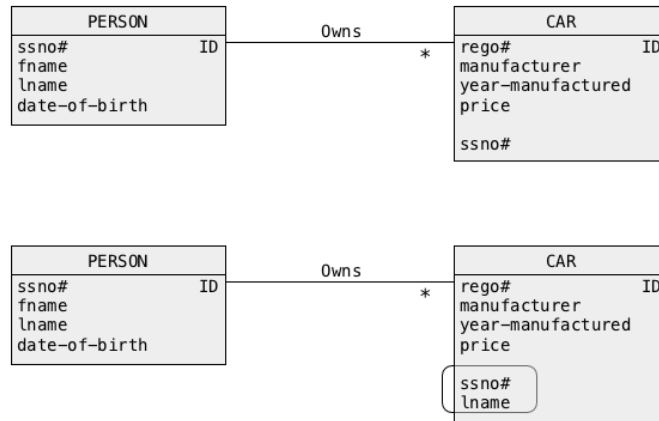
Simplifications

Migration of identifiers

Migration of attributes from “one” to “many” side

Migration of attributes from "one" to "many" side

It is always possible to copy an attribute from "one" side of an association to "many" side (and NOT the opposite) of the association



Performance related observations:

Copying nonID attributes allows for elimination of join operations in some queries


Copying nonID attributes destroys normalization and introduces redundancies

Copying nonID attributes is called as **denormalization**

Migration of attributes from "one" to "many" side

Applications

- Find the last names of people who own a car manufactured in year 2000
- Find how old are the cars owned by Smith



Performance related observations:


To process both queries there is no need to join the relational tables **PERSON** and **CAR**

A relational table **CAR** is extended by a redundant column **lname**

If a relational table **CAR** contains **n** rows then **n * size(lname)** extra bytes are required to implement the table

Migration of attributes from "one" to "many" side

What about normalization ?



Functional dependencies:

CAR:

rego# → manufacturer

...

rego# → ssno#

ssno# → lname

Minimal key: (**rego#**)

2NF because every nonprime attribute is fully functionally dependent on a minimal key

Migration of attributes from "one" to "many" side

What about normalization ?

Functional dependencies:

CAR:

$\text{rego\#} \rightarrow \text{manufacturer}$

...

$\text{rego\#} \rightarrow \text{ssno\#}$

$\text{ssno\#} \rightarrow \text{lname}$

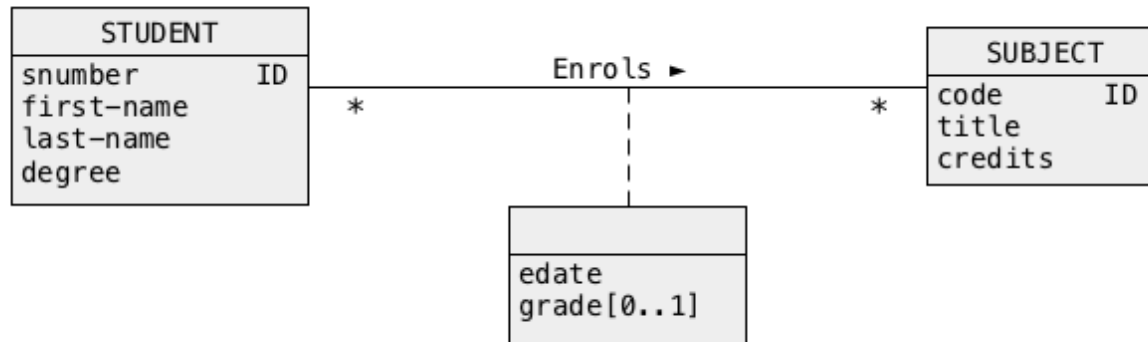
Minimal key: (rego\#)

Not 3NF because of transitive functional dependencies

$\text{rego\#} \rightarrow \text{ssno\#}$ and $\text{ssno\#} \rightarrow \text{lname}$

Migration of attributes from "one" to "many" side

Example




Assume that we would like to speed up processing of the following classes of queries:

- Find the titles of subjects enrolled by the students who belong to a given degree
- Find the names of degrees taken by the students who enrolled a subject with a give title

The attributes **title** and **degree** must be copied to the same class

Migration of attributes from "one" to "many" side

After simplification

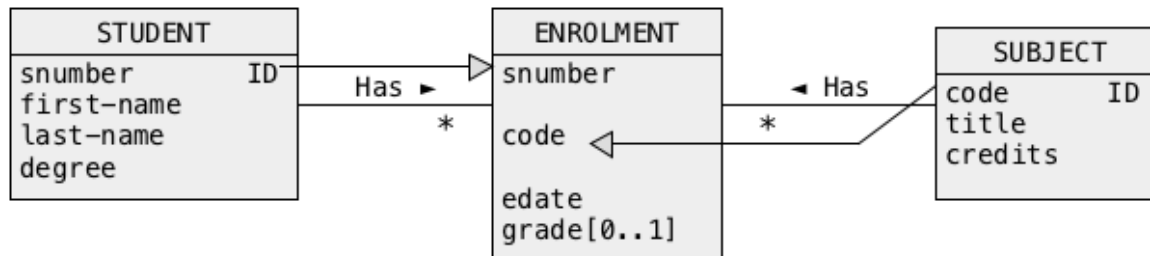


Many-to-many association described by the link attributes **edate** and **grade[0..1]** is replaced with a class **ENROLMENT** and "one-to-many" association **STUDENT** Has **ENROLMENT** and "one-to-many" association **SUBJECT** Has **ENROLMENT**

A class **ENROLMENT** is described by the attributes **edate** and **grade[0..1]**

Migration of attributes from "one" to "many" side

After implementation of associations




An identifier **snumber** of a class **STUDENT** is copied to a class **ENROLMENT**

An identifier **code** of a class **SUBJECT** is copied to a class **ENROLMENT**

Migration of attributes from "one" to "many" side

After denormalization



An attribute **degree** in a class **STUDENT** is copied to a class **ENROLMENT**

An attribute **title** in a class **SUBJECT** is copied to a class **ENROLMENT**

Now, the attributes **degree** and **title** can be used in the same query without the joins of relational tables

The design is redundant because the values of attributes **degree** and **title** are repeated each time a student enrolls a subject