

# CSCI317 Database Performance Tuning

## Introduction to Indexing

Dr Janusz R. Getta

School of Computing and Information Technology -  
University of Wollongong

# Introduction to Indexing

## Outline

### Index ? What is it ?

#### Index versus indexed file organization

#### Primary (unique) index

#### Secondary (nonunique) index

#### Clustered index

#### B\*-tree index implementation

#### Traversals of B\*-tree index

#### Hash-based index implementation

# Index ? What is it ?

An **index** is a data structure that organizes data records on persistent storage to optimize certain kinds of retrieval operations

An **index** is used to efficiently retrieve all records that satisfy a search condition on the search key fields of the index

An index is a function  $f: K \rightarrow \wp(id_R)$  where  $K$  is set of keys and

$\wp(id_R)$  is a powerset (a set of all sets) of identifiers (addresses)  $id_R$  of the records in a set  $R$

Let **EMP** be a relational table over a relational schema

**Employee(enumber, name, department)**

Then,  $F_{\text{department}}: \text{domain}(\text{department}) \rightarrow \wp(id_{\text{EMP}})$  is a function that maps the names of departments in  $\text{domain}(\text{DEPARTMENT})$  into the sets of identifiers of rows  $\wp(id_{\text{EMP}})$  in relational table **EMP**

$F_{\text{department}}(\text{'Sales'})$  returns the identifiers of all rows where a value of attribute **department** is equal to **'Sales'**

# Introduction to Indexing

## Outline

Index ? What is it ?

Index versus indexed file organization

Primary (unique) index

Secondary (nonunique) index

Clustered index

B\*-tree index implementation

Traversals of B\*-tree index

Hash-based index implementation

# Index versus indexed file organization

An **indexed file organization** (**index organized file**) is a function  $f : K \rightarrow \wp(R)$  where  $K$  is a set of keys and  $\wp(R)$  is a powerset (a set of sets) of records  $R$

An **index** maps a **value** into **set of row identifiers**

An **index organized file** maps a **value** into a **set of records**

A **relational table** can be **indexed** or it can be **index organized**

An **indexed relational table** consists of several **index(es)** created separately from implementation of a **relational table** itself

An **index organized relational table** consists only of implementation of one **index** where an **index key** is the same as a **relational schema** of an **index organized table**

Indexing in database systems is **transparent to data manipulation and data retrieval operations**

It means that a database system automatically **modifies an index** and automatically decides whether an **index is used for search**

# Introduction to Indexing

## Outline

Index ? What is it ?

Index versus indexed file organization

Primary (unique) index

Secondary (nonunique) index

Clustered index

B\*-tree index implementation

Traversals of B\*-tree index


Hash-based index implementation

# Primary (unique) index

A **primary (unique) index** is an index on a set of attributes equal to **primary** or **candidate key**

A **primary index** is a function  $f : K \rightarrow id_R$  where  $K$  is a set of **key values** and  $id_R$  is a set of identifiers (physical addresses) of rows in a relational table  $R$

A **primary index** maps and **index key** into a **single row identifier** (physical address of a row)



$F_{enum} : \text{domain(enum)} \rightarrow id_{EMPLOYEE}$

TOP

Created by Janusz R. Getta, CSCI317 Database Performance Tuning, SIM, Session 3, 2022

7/32

# Primary (unique) index

A **primary key** in a **relational table** is always **automatically indexed** by a database system

For example, a relational table **EMPLOYEE** created over a relational schema **Employee(enum, name, department)** where **enum** is a **primary key** has an index automatically created on an attribute (**enum**)

For example, a relational table **ENROLMENT** created over a relational schema **Enrolment(snumber, code, edate)** where (**snumber, code**) is a **primary key** has an index automatically created on a set of attributes (**snumber, code**)

An index on (**snumber, code**) is a **composite index**



# Introduction to Indexing

## Outline

Index ? What is it ?

Index versus indexed file organization

Primary (unique) index

Secondary (nonunique) index

Clustered index

B\*-tree index implementation

Traversals of B\*-tree index


Hash-based index implementation

# Secondary (nonunique) index

A **secondary index** is an index which is not **primary**

A **secondary index** is a function  $f: K \rightarrow \wp(R)$  where  $K$  is a set of **key values** and  $id_R$  is a set of identifiers (physical addresses) of rows in a relational table  $R$

A **secondary index** maps an **index key** into a **single row identifier** (physical address of a row)



$$F_{\text{department}}: \text{domain}(\text{department}) \rightarrow \wp(\text{id}_{\text{EMPLOYEE}})$$

# Secondary (nonunique) index

For example, an index on an attribute (**name**) in a relational table **EMPLOYEE** created over a relational schema **Employee(enum, name, department)** is a **secondary (nonunique) index**

For example, an index on a set of attributes (**name, department**) in a relational table **EMPLOYEE** created over a relational schema **Employee(enum, name, department)** is a **secondary index**

For example, an index on an attribute (**snumber**) in a relational table **ENROLMENT** created over a relational schema **Enrolment(snumber, code, edate)** is a **secondary index**

An index on a set of attributes (**enum, name**) in a relational table **EMPLOYEE** created over a relational schema **Employee(enum, name, department)** is still a primary index because (**enum, name**) is a superkey

# Introduction to Indexing

## Outline

Index ? What is it ?

Index versus indexed file organization

Primary (unique) index

Secondary (nonunique) index

Clustered index

B\*-tree index implementation


Traversals of B\*-tree index

Hash-based index implementation

# Clustered index

A **clustered index** is an index organized such that the ordering of rows is the same as ordering of keys in the index

A clustered index is a function  $f: K \rightarrow id_R$  where  $K$  is a set of keys and  $id_R$  is a set of row identifiers (addresses) in a relational table  $R$  such that  $f(v)$  returns row identifier (address) of the first row in a sequence of rows such that a value of attribute  $K$  is equal to  $v$



$f_{\text{department}}: \text{domain}(\text{department}) \rightarrow id_{\text{EMPLOYEE}}$

# Clustered index

Every **primary index** is **clustered**

**Clustered index** provides faster access to data than **nonclustered secondary** index

**Clustered index** has a very negative impact on performance of **INSERT** and **UPDATE** SQL statements

Therefore, **clustered indexing** should be applied to mainly to **read-only data**

# Introduction to Indexing

## Outline

Index ? What is it ?

Index versus indexed file organization

Primary (unique) index

Secondary (nonunique) index

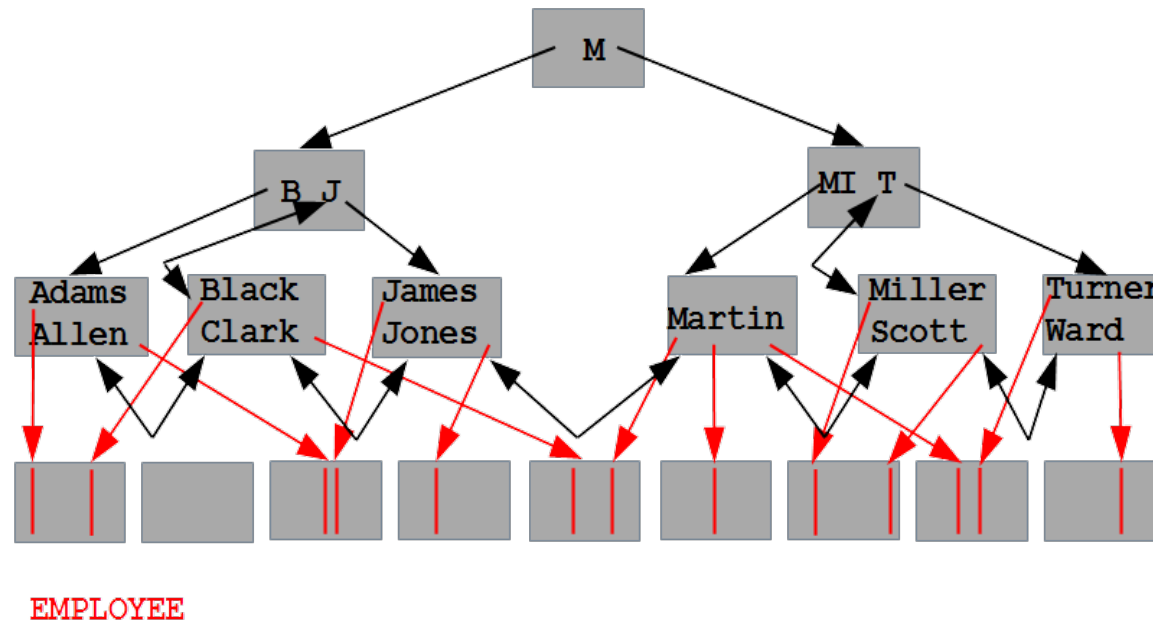
Clustered index

B\*-tree index implementation

Traversals of B\*-tree index

Hash-based index implementation

# B\*-tree index implementation



B\*-tree can be traversed either:

- vertically from root to leaf level of a tree
- horizontally either from left corner of leaf level to right corner of leaf level or the opposite
- vertically and later on horizontally either towards left lower corner or right lower corner of leaf level



# B\*-tree index implementation

The **height** of **B\*-tree** is a length of path from root node to any of leaf nodes measured either in a number of nodes or edges along the path

The **fanout** of **B\*-tree** is an average number of children nodes attached to a non-leaf node

If **B\*-tree** has the **height** equal to  $h$  and the **fanout** equal to  $f$  then such a tree has  $f^h$  leaf nodes

If  $k$  is the total number of keys then  $k = f^h$

Hence  $\log_f(k) = h$

The **height** of a tree  $h = \log_f(k) + 1$

# Introduction to Indexing

## Outline

Index ? What is it ?

Index versus indexed file organization

Primary (unique) index

Secondary (nonunique) index

Clustered index

B\*-tree index implementation


Traversals of B\*-tree index

Hash-based index implementation

# Traversals of B\*-tree index

## Vertical traversal

- **Vertical traversal** passes through B\*-tree from the root node to one of the leaf nodes



- The total number of data blocks read during a vertical scan of B\*-tree is equal to the height of B\*-tree  $h$  measured in the number of nonleaf nodes  $+ 1$
- The total number of data blocks read during a vertical traversal of B\*-tree is equal to  $\log_f(k) + 1$  where  $f$  is a fanout and  $k$  is the total number of keys included in the index

# Traversals of B\*-tree index

An index on a primary key (`enum`) in a relational table `EMPLOYEE` created over a relational schema `Employee(enum, name, department, salary)` is always built automatically by a database system

A name of an index is the same as a name of primary key constraint in a relational table `EMPLOYEE`

The following queries are processed through a **vertical traversal** of an index on (`enum`)

```
SELECT *  
FROM EMPLOYEE  
WHERE enum = 007;
```

[SQL](#)

```
SELECT *  
FROM EMPLOYEE  
WHERE enum = 007 AND department = 'MI6';
```

[SQL](#)

```
SELECT enum  
FROM EMPLOYEE  
WHERE enum = 007;
```

[SQL](#)

# Traversals of B\*-tree index

Assume that we created an index on attribute (**name**) in a relational table **EMPLOYEE** created over a relational schema

**Employee(enum, name, department, salary)**

The following queries will be processed through a **vertical traversal** of an index on (**name**)

```
SELECT *  
FROM EMPLOYEE  
WHERE name = 'James';
```

[SQL](#)

```
SELECT *  
FROM EMPLOYEE  
WHERE name = 'James' and department = 'MI6';
```

[SQL](#)

```
SELECT count(*)  
FROM EMPLOYEE  
WHERE name = 'James'
```

[SQL](#)

# Traversals of B\*-tree index

Assume that we created an index on the attributes  
(**name**, **department**) in a relational table **EMPLOYEE** created over a  
relational schema **Employee(enum, name, department, salary)**

The following queries are processed through a **vertical traversal** of an  
index on (**name**, **department**)

```
SELECT *  
FROM EMPLOYEE  
WHERE name = 'James' and department = 'MI6';
```

[SQL](#)

```
SELECT count(*)  
FROM EMPLOYEE  
WHERE name = 'James' and department = 'MI6';
```

[SQL](#)


```
SELECT *  
FROM EMPLOYEE  
WHERE name = 'James' and department = 'MI6' and salary > 1000;
```

[SQL](#)

# Traversals of B\*-tree index

## Horizontal traversal

- **Horizontal traversal** passes through B\*-tree from the leftmost (rightmost) leaf level node to the rightmost (leftmost) leaf level node



- The total number of data blocks read during a **horizontal traversal** of B\*-tree is equal to the total number of leaf level nodes **w**
- The **height** of B\*-Tree and the **total number of leaf level nodes** may be obtained from a data dictionary view **USER\_TABLES** after processing of **ANALYZE TABLE** statement

# Traversals of B\*-tree index

The following queries are processed through a **horizontal traversal** of leaf level of an index on (**enum**)

```
SELECT COUNT(*)  
FROM EMPLOYEE;
```

[SQL](#)

```
SELECT COUNT(enum)  
FROM EMPLOYEE;
```

[SQL](#)

```
SELECT COUNT(name) /* Only if name IS NOT NULL */  
FROM EMPLOYEE;
```

[SQL](#)

```
SELECT enum  
FROM EMPLOYEE;
```

[SQL](#)

```
SELECT enum, COUNT(*)  
FROM EMPLOYEE  
GROUP BY enum;
```

[SQL](#)

```
SELECT enum  
FROM EMPLOYEE  
ORDER BY enum;
```


[SQL](#)



# Traversals of B\*-tree index

## Vertical and horizontal traversal

- **Vertical and horizontal traversal** passes through B\*-tree from the root node to one of the leaf nodes and later on the leaf nodes towards the rightmost (leftmost) leaf node



- The total number of data blocks read during a vertical and horizontal traversal of B\*-tree is equal to  $\log_f(k) + 1 + p \cdot w$  where  $f$  is a fanout,  $k$  is the total number of keys,  $w$  is the total number of leaf level blocks and  $p$  is a fraction of  $w$  to be traversed

# Traversals of B\*-tree index

The following queries can be processed through a vertical traversal and later on horizontal traversal of an index on (enum)

```
SELECT *  
FROM EMPLOYEE  
WHERE enum > 300;
```

[SQL](#)

```
SELECT count(*)  
FROM EMPLOYEE  
WHERE enum < 007;
```

[SQL](#)

```
SELECT *  
FROM EMPLOYEE  
WHERE enum > 300 and salary > 1000;
```

[SQL](#)

# Traversals of B\*-tree index

Assume that we created an index on the attributes (name, department) in a relational table EMPLOYEE created over a relational schema Employee(enum, name, department, salary)

The following queries can be processed through a vertical traversal and later on horizontal traversal of an index on (name, department)

```
SELECT *  
FROM EMPLOYEE  
WHERE name > 'James';
```

[SQL](#)

```
SELECT count(*)  
FROM EMPLOYEE  
WHERE name <= 'James';
```

[SQL](#)

```
SELECT *  
FROM EMPLOYEE  
WHERE name = 'James' and department > 'MI6';
```

[SQL](#)

# Traversals of B\*-tree index

Assume that we created a **composite key index** on the attributes (**name**, **department**) in a relational table **EMPLOYEE** created over a relational schema **Employee(enum, name, department, salary)**

The following queries **can be processed** through a **vertical traversal** and later on **horizontal traversal** of an index on (**name**, **department**)

```
SELECT *  
FROM EMPLOYEE  
WHERE name > 'James' and salary > 1000;
```

[SQL](#)

```
SELECT name, count(*)  
FROM EMPLOYEE  
WHERE name > 'James' and salary > 1000  
GROUP BY name;
```

[SQL](#)

```
SELECT *  
FROM EMPLOYEE  
WHERE name > 'James' and salary > 1000  
ORDER BY name;
```

[SQL](#)

# Introduction to Indexing

## Outline

Index ? What is it ?

Index versus indexed file organization

Primary (unique) index

Secondary (nonunique) index

Clustered index

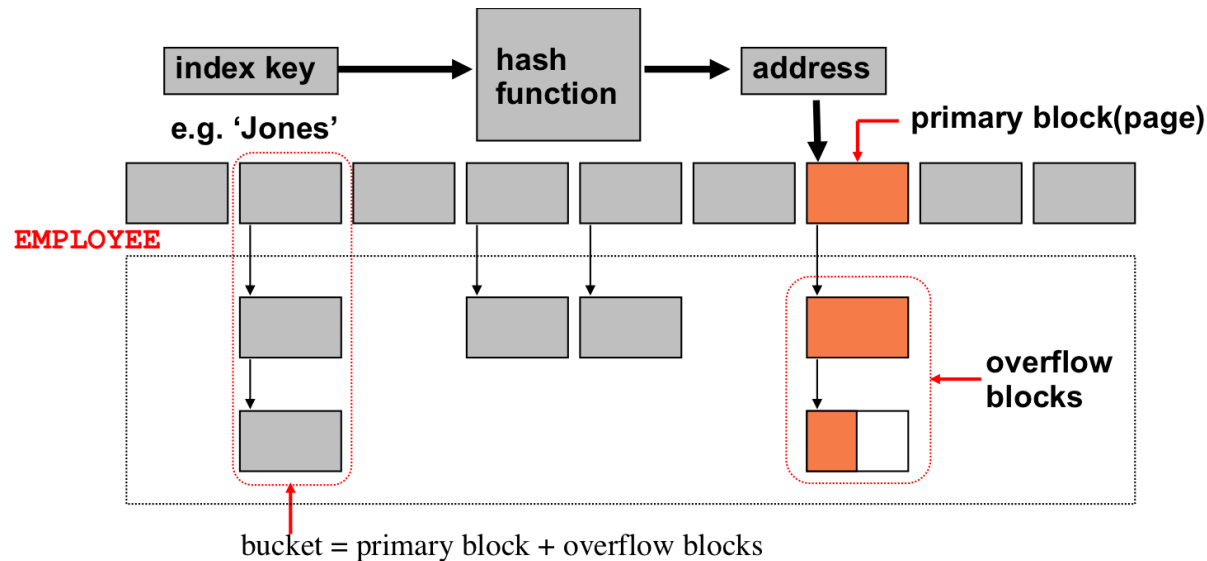
B\*-tree index implementation

Traversals of B\*-tree index

Hash-based index implementation

# Hash-based index implementation

A **hash-based index** consists of a sequence of **primary** data blocks and the "**buckets**" of **overflow** data blocks linked to the **primary** data blocks



To insert a row into **hash-based index** a **hash function** is computed over a **key**

A result of **hash function** determines a **primary block** or a "**bucket**" of **overflow** blocks that should be used for insertion

# Hash-based index implementation

A **hash function** has the following properties

- A **hash function**  $h$  does not have an inverse function, for some keys  $k_1, k_2, \dots, k_n$   $h(k_1) = h(k_2) = \dots = h(k_n)$
- A result of **hash function** determines a **primary block** or a "bucket" of **overflow** blocks that should be used for insertion
- Quality of hashing depends on how evenly a hash function distributes the records over buckets
- The best quality is achieved if size of each bucket is more or less the same
- A sample hash function: (folding and xor-ing  $n$  bits ) mod number-of-buckets

# References

Elmasri R. and Navathe S. B., Fundamentals of Database Systems,  
Chapter 17 Indexing Structures for Files and Physical Database Design,  
7th ed., The Person Education Ltd, 2017