

CSCI317 Database Performance Tuning

Index Selection Guidelines

Dr Janusz R. Getta

School of Computing and Information Technology -
University of Wollongong

Index Selection Guidelines

Outline

To use index or not to use index ?

Efficiency of indexing

A wise choice of an index key

Cluster index versus non-cluster index

Hash index versus B*-tree index

Balancing the costs of index maintenance

To use an index or not to use an index ?

Main principle

- Do not build index unless some query (including the query components of updates and deletions) benefit from it

Selectivity of an attribute **A** in a relational table **T** is defined as a result of a query

```
SELECT COUNT(DISTINCT A) / COUNT(*)  
FROM T;
```

Selectivity of attribute A

Selectivity of an attribute is a number in a range $(0, 1]$

Primary key and candidate key have the highest selectivity equal to **1**

Selectivity of a set of attributes (A, B, \dots) is computed in the same way a selectivity of a single attribute

Index Selection Guidelines

Outline

To use index or not to use index ?

Efficiency of indexing


A wise choice of an index key

Cluster index versus non-cluster index

Hash index versus B*-tree index

Balancing the costs of index maintenance

Efficiency of indexing



Index Selection Guidelines

Outline

To use index or not to use index ?

Efficiency of indexing

A wise choice of an index key

Cluster index versus non-cluster index

Hash index versus B*-tree index

Balancing the costs of index maintenance

A wise choice of an index key

An exact-match selection condition such as `attribute=value` suggests that we should consider an index on `attribute` when selectivity of `attribute` is high

If `attribute` is frequently updated or new values are added and old values are deleted then it should be **B-Tree** based index (majority of the cases)

If `attribute` does not change frequently then we should consider **hash-based** index

If a relational table that includes `attribute` changes frequently then we must use **non-cluster based** index

If a relational table that includes `attribute` has very infrequent changes then **cluster based** index is a good option

A wise choice of an index key

A range selection condition such as `attribute > value` means that we must use **B-Tree based** index on `attribute`

If a relational table that includes `attribute` changes frequently then we must use **non-cluster** based index

If a relational table that includes `attribute` has very infrequent changes then **cluster based** index is a good option

A wise choice of an index key

Index with a **composite key** should be considered when **WHERE** clause includes a condition on more than one attribute and a condition is a conjunction (**AND**) of elementary terms

For example a query like

```
SELECT *  
FROM EMPLOYEE  
WHERE fname = 'James' AND lname = 'Bond';
```

Benefits from a composite index key

benefits of from a composite key index (**fname, lname**)

An order of attributes in a composite key index is very important

For example a query like

```
SELECT *  
FROM EMPLOYEE  
WHERE lname = 'Bond';
```

May benefit from a composite index key

benefits of from a composite key index (**lname, fname**)

A wise choice of an index key

It is always worth to consider an index that allows for **index only** processing of a query, for example a query like

```
SELECT fname, lname  
FROM EMPLOYEE
```

Index only processing of a query

does not need access to a relational table **EMPLOYEE** and entire processing can be done on a composite key index (**lname, fname**)

Another example of **index only** processing of a query

```
SELECT city  
FROM EMPLOYEE  
WHERE city LIKE 'A%';
```

Index only processing of a query

does not need access to a relational table **EMPLOYEE** and entire processing can be done on a single attribute index (**city**)

A wise choice of an index key

An index can be used for aggregation, sorting, and grouping, for example, query like

```
SELECT COUNT(*)  
FROM EMPLOYEE;
```

Horizontal traversal of an index

does not need access to a relational table **EMPLOYEE** and entire processing can be done on an index on primary key

For example, a query like

```
SELECT *  
FROM EMPLOYEE  
ORDER BY lname
```

Horizontal traversal of an index

does not need to sort a relational table **EMPLOYEE** and sorting can be replaced with horizontal traversal through leaf level of a composite key index (**lname, fname**)

A wise choice of an index key

An index can be used for aggregation, sorting, and grouping, for example, a query like

```
SELECT lname, COUNT(*)  
FROM EMPLOYEE  
GROUP BY lname;
```

Horizontal traversal of an index

does not need to group a relational table **EMPLOYEE** over **lname** and entire processing can be done by a horizontal scan through leaf level of a composite key index (**lname, fname**)

Index Selection Guidelines

Outline

To use index or not to use index ?

Efficiency of indexing

A wise choice of an index key

Cluster index versus non-cluster index

Hash index versus B*-tree index

Balancing the costs of index maintenance

Cluster index versus non-cluster index

Clustering has a very positive impact on performance of query processing

Clustering has a very negative impact on performance on insert/update/delete operations

Range queries are likely to benefit from **clustering**

If an index enables index-only query computation strategy then such index does not need to be **cluster index**

Index Selection Guidelines

Outline

To use index or not to use index ?

Efficiency of indexing

A wise choice of an index key

Cluster index versus non-cluster index

Hash index versus B*-tree index

Balancing the costs of index maintenance

Hash index versus B*-Tree index

B*-Tree index is preferable when a table is frequently updated

Only B*-Tree index can be used for range queries

Hash index is preferable for equality queries on read-only relational table

Hash index better supports hash-based implementation of join operation

Index Selection Guidelines

Outline

To use index or not to use index ?

Efficiency of indexing

A wise choice of an index key

Cluster index versus non-cluster index

Hash index versus B*-tree index

Balancing the costs of index maintenance

Balancing the costs of index maintenance

If **index maintenance** slows down important **UPDATE/DELETE/INSERT** operations - **drop an index**

Index may speed up **UPDATE/DELETE/INSERT** operations when their computation involves query processing

References

[Cookbook, How to measure and how to improve performance of database applications, how to choose the best index, how to analyze index structures ?](#)

Ramakrishnan R., Gehrke J., Database Management Systems, chapters 20.1-20.3

Lightstone, S., Teorey T., Nadeau T., Physical Database Design, The Database Professional's Guide to Exploiting Indexes, Views, Storage, and More, Morgan Kaufmann Publishers, 2007, chapter 4

Sasha, D., Bonnet, P., Database Tuning Principles, Experiments, and Troubleshooting Techniques, Morgan Kaufmann, 2003, chapter 3

Kyte, T., Expert Oracle Database Architecture, 9i and 10g Programming Techniques and Solutions, APress, 2005, chapter 11