CSCI317 Database Performance Tuning

Indexing Examples

Dr Janusz R. Getta

School of Computing and Information Technology -
University of Wollongong

Indexing Examples

Outline

Examples

Examples

Consider a relational table
`PRODUCT(PNUM, PNAME, PRICE, MANUFACTURER)`

Assume that

- the table contains 10^5 rows,
- an attribute `PNUM` is a primary key,
- the products are manufactured by `100` manufacturers,
- the table contains `1000` product names,
- an average number of rows per disk block (blocking factor) is `50`
- a database administrator created a non-clustered B*-tree index on attribute `MANUFACTURER`,
- a fanout of B*-tree is equal to `f`
- a width (total number of blocks at leaf level) of B*-tree implementing an index on the primary key is equal to w_p
- a width (total number of blocks at leaf level) of B*-tree implementing an index on an attribute `MANUFACTURER` is equal to w_m

Example 1

To compute a query

Projection query, no index can be used

```
SELECT PNAME, PRICE
FROM  PRODUCT;
```

we have to read $10^5/50$ (`total number of rows/blocking factor`) blocks because the entire relational table `PRODUCT` must be read

A fragment of respective query processing plan

A fragment of query processing plan

```
-------------------------------------------- ...
| Id | Operation          | Name    |
-------------------------------------------- ...
| 0  | SELECT STATEMENT   |         |
| 1  |  TABLE ACCESS FULL | PRODUCT |
-------------------------------------------- ...
```

Example 2

To compute a query

Selection and projection query, index on primary key can be used

```
SELECT PNAME, PRICE
FROM  PRODUCT
WHERE PNUM = 12345;
```

we have to read $\lceil \log_f 10^5 + 1 \rceil + 1$ blocks because the primary key `PNUM` is automatically indexed

A fragment of respective query processing plan

A fragment of query processing plan

```
-------------------------------------------------------- ...
| Id | Operation                     | Name          |
-------------------------------------------------------- ...
| 0  | SELECT STATEMENT              |               |
| 1  |  TABLE ACCESS BY INDEX ROWID  | PRODUCT       |
|*2  |   INDEX UNIQUE SCAN           | PRODUCT_PKEY  |
-------------------------------------------------------- ...
```

Example 3

To compute a query

Selection and projection query, index on primary key can be used

```
SELECT PNAME, PRICE
FROM   PRODUCT
WHERE PNUM = 1234567 AND
      PRICE <= 20.0;
```

we have to read $[\log_f 10^5 + 1] + 1$ blocks because the primary key `PNUM` is automatically indexed

A fragment of respective query processing plan

A fragment of query processing plan

```
------------------------------------------------------- ...
| Id | Operation                     | Name          |
------------------------------------------------------- ...
| 0  | SELECT STATEMENT              |               |
| 1  |  TABLE ACCESS BY INDEX ROWID  | PRODUCT       |
|*2  |   INDEX UNIQUE SCAN           | PRODUCT_PKEY  |
------------------------------------------------------- ...
```

Example 4

To compute a query

Selection and projection query, index on primary key cannot be used due to disjunction (OR)

```
SELECT PNAME, PRICE
FROM   PRODUCT
WHERE  PNUM = 1234567 OR
       PRICE <= 20.0;
```

we have to read $10^5/50$ (`total number of rows/blocking factor`) blocks because the entire relational table `PRODUCT` must be read

A fragment of query processing plan

```
------------------------------------------- ...
| Id | Operation          | Name    |
------------------------------------------- ...
| 0  | SELECT STATEMENT   |         |
| 1  |  TABLE ACCESS FULL | PRODUCT |
------------------------------------------- ...
```

Example 5

To compute a query

```
                                    Selection and projection query, index on MANUFACTURER can be used
SELECT PNAME, PRICE
FROM  PRODUCT
WHERE MANUFACTURER = 'GoldenBolts';
```

we have to read
$[\log_f 10^2 + 1] + ((10^5/100) + 10^5/(100*50))/2$ data blocks

Index processing: $[\log_f 10^2 + 1]$ read block operations

Total number of row identifiers found: $10^5/100$

Table processing:

- The best case (full clustering: $10^5/(100*50)$ read block operations
- The worst case (each row in a different block): $10^5/100$ read block operations

Example 5 (continuation)

To compute a query

Selection and projection query, index on MANUFACTURER can be used

```
SELECT PNAME, PRICE
FROM  PRODUCT
WHERE MANUFACTURER = 'GoldenBolts';
```

we have to read

$[\log_f 10^2 + 1] + ((10^5/100) + 10^5/(100*50))/2$ data blocks

A fragment of query processing plan

```
------------------------------------------------------ ...
| Id | Operation                    | Name    |
------------------------------------------------------ ...
| 0  | SELECT STATEMENT             |         |
| 1  |  TABLE ACCESS BY INDEX ROWID | PRODUCT |
|*2  |   INDEX RANGE SCAN           | PIDX    |
------------------------------------------------------ ...
```

Example 6

To compute a query

Counting the rows, and index on primry key can be used

```
SELECT COUNT(*)
FROM PRODUCT;
```

we have to read w_p data blocks (width of leaf level on an index on primary key)

Index processing: w_p read block operations

Counting row identifiers: 0 read block operations, counting is performed in transient memory

A fragment of query processing plan

```
-------------------------------------------------- ...
| Id  | Operation             | Name         |
-------------------------------------------------- ...
| 0   | SELECT STATEMENT      |              |
| 1   |  SORT AGGREGATE       |              |
| 2   |   INDEX FAST FULL SCAN| PRODUCT_PKEY |
-------------------------------------------------- ...
```

Example 7

To compute a query

Counting the rows with selection, an index on MANUFACTURER can be used

```
SELECT COUNT(*)
FROM  PRODUCT;
WHERE MANUFACTURER = 'GoldenBolts';
```

we have to read $\lceil \log_f 10^2 + 1 \rceil$ data blocks

Index processing: $\lceil \log_f 10^2 + 1 \rceil$ read block operations

Counting row identifiers: 0 read block operations, counting is performed in transient memory

A fragment of query processing plan

```
------------------------------------------- ...
| Id | Operation          | Name   |
------------------------------------------- ...
| 0  | SELECT STATEMENT   |        |
| 1  |  SORT AGGREGATE    |        |
|*2  |   INDEX RANGE SCAN | PIDX   |
------------------------------------------- ...
```

Example 8

To compute a query

Projection and selection with patter matching

```
SELECT MANUFACTURER
FROM   PRODUCT
WHERE MANUFACTURER LIKE 'M%';
```

we have to read w_m data blocks (width of a leaf level of an index on MANUFACTURER)

Index processing: w_m read block operations

A fragment of query processing plan

```
-------------------------------------------- ...
| Id | Operation          | Name |
-------------------------------------------- ...
| 0  | SELECT STATEMENT   |      |
|*1  |  INDEX RANGE SCAN | PIDX |
-------------------------------------------- ...
```

Example 9

To compute a query

```
                Projection query with elimination of duplicates, an index on MANUFACTURER can be used
SELECT DISTINCT MANUFACTURER
FROM  PRODUCT;
```

we have to read w_m data blocks

Index processing: w_m read block operations (horizontal traversal through leaf level on an index on MANUFACTURER)

```
                                                A fragment of query processing plan
------------------------------------------------ ...
| Id  | Operation             | Name          |
------------------------------------------------ ...
| 0   | SELECT STATEMENT      |               |
| 1   |  SORT UNIQUE          |               |
| 2   |   INDEX FAST FULL SCAN| PRODUCT_PKEY  |
------------------------------------------------ ...
```

Example 10

To compute a query

Aggregation query, an index on MANUFACTURER can be used

```
SELECT MANUFACTURER, COUNT(*)
FROM PRODUCT
GROUP BY MANUFACTURER;
```

we have to read w_m data blocks

Index processing: w_m read block operations

A fragment of query processing plan

```
------------------------------------------- ...
| Id | Operation             | Name         |
------------------------------------------- ...
| 0  | SELECT STATEMENT      |              |
| 1  |  SORT GROUP BY NOSORT |              |
| 2  |   INDEX FULL SCAN     | PRODUCT_PKEY |
------------------------------------------- ...
```

Examples (continuation)

Consider a relational table

`PRODUCT(PNUM, PNAME, PRICE, MANUFACTURER)`

Assume that

- a database administrator created a non-clustered B*-tree index on attribute `PNAME`,
- the table contains `1000` distinct product name
- a fanout of B*-tree is equal to `f`
- a width (total number of blocks at leaf level) of B*-tree implementing an index on an attribute `PNAME` is equal to w_n

Example 10

To compute a query

```
  Selection query with conjunction WHERE condition, both indexes on MANUFACTURER and PNAME can be used
SELECT *
FROM PRODUCT
WHERE PNAME = 'bolt' AND
      MANUFACTURER = 'Golden Bolts';
```

we have to read $[\log_f 1000 + 1] + [\log_f 100 + 1] + 1$ data blocks

Index processing: $[\log_f 1000 + 1]$ read block operations to process an index on `PNAME`

The total number of row identifiers found: $10^5/1000$

Index processing: $[\log_f 100 + 1]$ read block operations to process an index on `MANUFACTURER`

The total number of row identifiers found: $10^5/100$

Intersection of the sets of row identifiers obtained from index processing: 0 read block operations

Example 10

Total number of row identifiers found: $10^5/(1000*100) = 1$

Table processing: 1 read block operations

Example 11

To compute a query

```
  Selection query with disjunction WHERE condition, both indexes on MANUFACTURER and PNAME can be used
SELECT *
FROM PRODUCT
WHERE PNAME = 'bolt' OR
      MANUFACTURER = 'Golden Bolts';
```

we have to read $[\log_f 1000 + 1] + [\log_f 100 + 1] + ((10^2 + 10^3 - 1) + (10^2 + 10^3 - 1)/50)/2$ data blocks

Index processing: $[\log_f 1000 + 1]$ read block operations to process an index on PNAME

The total number of row identifiers found: $10^5/1000$

Index processing: $[\log_f 100 + 1]$ read block operations to process an index on MANUFACTURER

The total number of row identifiers found: $10^5/100$

Example 11

Union of the sets of row identifiers obtained from index processing: 0 read block operations

Total number of row identifiers found: $(10^5/1000) + (10^5/100) - (10^5/(1000*100)) = 10^2 + 10^3 - 1$

Table processing: $((10^2 + 10^3 - 1) + (10^2 + 10^3 - 1)/50)/2$ read block operations

Example 12

To compute a query

```
Selection query with a range condition, an index on PNAME can be used
SELECT  *
FROM PRODUCT
WHERE PNAME > 'bolt';
```

we have to read $[\log_f 1000 + 1] + w_p * n/10^5 + (n + n/50)/2$ data blocks

Vertical index processing: $[\log_f 1000 + 1]$ read block operations to process an index on `PNAME`

Horizontal index processing: $w_p * n/10^5$ where w_p is the total number of leaf level blocks and n is the total number of rows that satisfy a condition `PNAME > 'bolt'`

References

Cookbook, How to measure and how to improve performance of database applications, how to choose the best index, how to analyze index structures ?

Ramakrishnan R., J. Gehrke Database Management Systems, chapters 8.1-8.3

Lightstone, S., Teorey T., Nadeau T., Physical Database Design, The Database Professional's Guide to Exploiting

Indexes, Views, Storage, and More, Morgan Kaufmann Publishers, 2007, chapter 4