

CSCI317 Database Performance Tuning

SQL Tuning (2)

Dr Janusz R. Getta

School of Computing and Information Technology -
University of Wollongong

SQL Tuning (2)

Outline

Index (nested loop) joins

Hash joins

Sort-merge joins

Order of join arguments in 2-way joins

Order of join arguments in m-way joins

Comparison of join algorithms

Optimization of joins through parameters

Outer joins

Nested queries

Index (nested loop) joins

Typical **index based join** uses an index on a **primary key**

Index based join

```
SELECT /*+ LEADING(ORDERS)
        USE_NL_WITH_INDEX(CUSTOMER CUSTOMER_PKEY) */ *
FROM ORDERS JOIN CUSTOMER
      ON ORDERS.O_CUSTKEY = CUSTOMER.C_CUSTKEY;
```

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
|-----|-----------------------------|---------------|------|-------|-------------|----------|
| 0 | SELECT STATEMENT | | 450K | 115M | 452K (1) | 00:00:18 |
| 1 | NESTED LOOPS | | | | | |
| 2 | NESTED LOOPS | | 450K | 115M | 452K (1) | 00:00:18 |
| * 3 | TABLE ACCESS FULL | ORDERS | 450K | 46M | 1950 (1) | 00:00:01 |
| * 4 | INDEX UNIQUE SCAN | CUSTOMER_PKEY | 1 | | 0 (0) | 00:00:01 |
| 5 | TABLE ACCESS BY INDEX ROWID | CUSTOMER | 1 | 159 | 1 (0) | 00:00:01 |

3 - filter("ORDERS"."O_CUSTKEY">=0)

4 - access("ORDERS"."O_CUSTKEY"="CUSTOMER"."C_CUSTKEY")

Index (nested loop) joins

If attributes from one table are only needed and index on a **foreign key** speeds up **index based join**

Creating and index on a foreign key

```
CREATE INDEX IDX ON ORDERS(O_CUSTKEY);
SELECT /*+ LEADING(ORDERS)
        USE_NL_WITH_INDEX(CUSTOMER CUSTOMER_PKEY) */ C_CUSTKEY, C_NAME
FROM ORDERS JOIN CUSTOMER
        ON ORDERS.O_CUSTKEY = CUSTOMER.C_CUSTKEY;
```

| Id | Operation | Name | Rows | Bytes | Cost | (%CPU) | Time |
|-----|-----------------------------|---------------|------|-------|------|--------|----------|
| 0 | SELECT STATEMENT | | 450K | 12M | 450K | (1) | 00:00:18 |
| 1 | NESTED LOOPS | | | | | | |
| 2 | NESTED LOOPS | | 450K | 12M | 450K | (1) | 00:00:18 |
| * 3 | INDEX FAST FULL SCAN | IDX | 450K | 2197K | 272 | (1) | 00:00:01 |
| * 4 | INDEX UNIQUE SCAN | CUSTOMER_PKEY | 1 | | 0 | (0) | 00:00:01 |
| 5 | TABLE ACCESS BY INDEX ROWID | CUSTOMER | 1 | 24 | 1 | (0) | 00:00:01 |

3 - filter("ORDERS"."O_CUSTKEY">=0)

4 - access("ORDERS"."O_CUSTKEY"="CUSTOMER"."C_CUSTKEY")

Index (nested loop) joins

If join is performed over **PK-FK** and all attributes selected are from **FK** side then there is no need for join

SELECT statement with redundant JOIN operation

```
SELECT /*+ LEADING(ORDERS)
        USE_NL_WITH_INDEX(CUSTOMER CUSTOMER_PKEY) */
        ORDERS.O_ORDERKEY, ORDERS.O_ORDERSTATUS
FROM ORDERS JOIN CUSTOMER
        ON ORDERS.O_CUSTKEY = CUSTOMER.C_CUSTKEY;
```

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
|----|-------------------|--------|------|-------|-------------|----------|
| 0 | SELECT STATEMENT | | 450K | 3515K | 1948 (1) | 00:00:01 |
| 1 | TABLE ACCESS FULL | ORDERS | 450K | 3515K | 1948 (1) | 00:00:01 |

SQL Tuning (2)

Outline

Index (nested loop) joins

Hash joins

Sort-merge joins

Order of join arguments in 2-way joins

Order of join arguments in m-way joins

Comparison of join algorithms

Optimization of joins through parameters

Outer joins

Nested queries

Hash joins

Hash join is the most effective implementation of **join operation** when two large relational tables are joined over an equality conditions

SELECT statement with JOIN operation

```
SELECT ORDERS.O_ORDERKEY, ORDERS.O_ORDERSTATUS, CUSTOMER.C_NAME
FROM ORDERS JOIN CUSTOMER
      ON ORDERS.O_CUSTKEY = CUSTOMER.C_CUSTKEY;
```

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|-----|-------------------|----------|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 450K | 15M | | 2843 (1) | 00:00:01 |
| * 1 | HASH JOIN | | 450K | 15M | 1584K | 2843 (1) | 00:00:01 |
| 2 | TABLE ACCESS FULL | CUSTOMER | 45000 | 1054K | | 282 (0) | 00:00:01 |
| * 3 | TABLE ACCESS FULL | ORDERS | 450K | 5712K | | 1949 (1) | 00:00:01 |

1 - access("ORDERS"."O_CUSTKEY"="CUSTOMER"."C_CUSTKEY")

3 - filter("ORDERS"."O_CUSTKEY">=0)

Hash joins

Changing an order of arguments in hash join may have an impact on performance

```
SELECT /*+ LEADING(ORDERS) */
```

SELECT statement with LEADING hint and JOIN operation

```
ORDERS.O_ORDERKEY, ORDERS.O_ORDERSTATUS, CUSTOMER.C_NAME
```

```
FROM ORDERS JOIN CUSTOMER
```

```
ON ORDERS.O_CUSTKEY = CUSTOMER.C_CUSTKEY;
```

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost | (%CPU) | Time |
|-----|-------------------|----------|-------|-------|---------|------|--------|----------|
| 0 | SELECT STATEMENT | | 450K | 15M | | 2844 | (1) | 00:00:01 |
| * 1 | HASH JOIN | | 450K | 15M | 10M | 2844 | (1) | 00:00:01 |
| * 2 | TABLE ACCESS FULL | ORDERS | 450K | 5712K | | 1949 | (1) | 00:00:01 |
| 3 | TABLE ACCESS FULL | CUSTOMER | 45000 | 1054K | | 282 | (0) | 00:00:01 |

1 - access("ORDERS"."O_CUSTKEY"="CUSTOMER"."C_CUSTKEY")

2 - filter("ORDERS"."O_CUSTKEY">=0)

Hash joins

Indexing can be used to speed up hash join when only the attributes from **PK** side are selected

Creating an index

```
CREATE INDEX IDX ON ORDERS(O_CUSTKEY);
```

SELECT statement with JOIN operation

```
SELECT CUSTOMER.C_CUSTKEY, CUSTOMER.C_NAME
FROM ORDERS JOIN CUSTOMER
      ON ORDERS.O_CUSTKEY = CUSTOMER.C_CUSTKEY;
```

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|-----|----------------------|----------|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 450K | 11M | | 970 (1) | 00:00:01 |
| * 1 | HASH JOIN | | 450K | 11M | 1496K | 970 (1) | 00:00:01 |
| 2 | TABLE ACCESS FULL | CUSTOMER | 45000 | 966K | | 282 (0) | 00:00:01 |
| * 3 | INDEX FAST FULL SCAN | IDX | 450K | 1757K | | 272 (1) | 00:00:01 |

```
1 - access("ORDERS"."O_CUSTKEY"="CUSTOMER"."C_CUSTKEY")
3 - filter("ORDERS"."O_CUSTKEY">=0)
```

Hash joins

Hash join cannot be used when the relational tables are joined over non equality join conditions

SELECT statement with JOIN operation

```
SELECT /*+ LEADING(ORDERS) USE_HASH(CUSTOMER) */ *
FROM ORDERS JOIN CUSTOMER
      ON ORDERS.O_TOTALPRICE > CUSTOMER.C_ACCTBAL;
```

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
|-----|-------------------|----------|-------|-------|-------------|----------|
| 0 | SELECT STATEMENT | | 20G | 5015G | 126M (1) | 01:22:09 |
| 1 | NESTED LOOPS | | 20G | 5015G | 126M (1) | 01:22:09 |
| 2 | TABLE ACCESS FULL | ORDERS | 450K | 46M | 1950 (1) | 00:00:01 |
| * 3 | TABLE ACCESS FULL | CUSTOMER | 44658 | 6934K | 280 (0) | 00:00:01 |

3 – filter("ORDERS"."O_TOTALPRICE">"CUSTOMER"."C_ACCTBAL")

SQL Tuning (2)

Outline

Index (nested loop) joins

Hash joins

Sort-merge joins

Order of join arguments in 2-way joins

Order of join arguments in m-way joins

Comparison of join algorithms

Optimization of joins through parameters

Outer joins

Nested queries

Sort-merge joins

Sort-merge join is used when the relational tables are joined over non equality join conditions

SELECT statement with JOIN operation

```
SELECT *
FROM ORDERS JOIN CUSTOMER
      ON ORDERS.O_TOTALPRICE > CUSTOMER.C_ACCTBAL;
```

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|-----|-------------------|----------|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 20G | 5015G | | 66173 (78) | 00:00:03 |
| 1 | MERGE JOIN | | 20G | 5015G | | 66173 (78) | 00:00:03 |
| 2 | SORT JOIN | | 45000 | 6987K | 15M | 1858 (1) | 00:00:01 |
| 3 | TABLE ACCESS FULL | CUSTOMER | 45000 | 6987K | | 283 (1) | 00:00:01 |
| * 4 | SORT JOIN | | 450K | 46M | 115M | 13009 (1) | 00:00:01 |
| 5 | TABLE ACCESS FULL | ORDERS | 450K | 46M | | 1950 (1) | 00:00:01 |

```
4 - access("ORDERS"."O_TOTALPRICE">"CUSTOMER"."C_ACCTBAL")
    filter("ORDERS"."O_TOTALPRICE">"CUSTOMER"."C_ACCTBAL")
```

Sort-merge joins

Sort-merge join can be enforced through a [hint](#)

SELECT statement with JOIN operation and a hint

```
SELECT /*+ USE_MERGE(ORDERS CUSTOMER) */
      ORDERS.O_ORDERKEY, ORDERS.O_ORDERSTATUS, CUSTOMER.C_NAME
FROM ORDERS JOIN CUSTOMER
      ON ORDERS.O_CUSTKEY = CUSTOMER.C_CUSTKEY;
```

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost | (%CPU) | Time |
|-----|-------------------|----------|-------|-------|---------|------|--------|----------|
| 0 | SELECT STATEMENT | | 450K | 15M | | 4685 | (1) | 00:00:01 |
| 1 | MERGE JOIN | | 450K | 15M | | 4685 | (1) | 00:00:01 |
| 2 | SORT JOIN | | 45000 | 1054K | 2840K | 599 | (1) | 00:00:01 |
| 3 | TABLE ACCESS FULL | CUSTOMER | 45000 | 1054K | | 282 | (0) | 00:00:01 |
| * 4 | SORT JOIN | | 450K | 5712K | 20M | 4086 | (1) | 00:00:01 |
| * 5 | TABLE ACCESS FULL | ORDERS | 450K | 5712K | | 1949 | (1) | 00:00:01 |

```
4 - access("ORDERS"."O_CUSTKEY"="CUSTOMER"."C_CUSTKEY")
filter("ORDERS"."O_CUSTKEY"="CUSTOMER"."C_CUSTKEY")
5 - filter("ORDERS"."O_CUSTKEY">=0)
```

Sort-merge joins

Sort-merge join can be improved through indexing

Creating an index

```
CREATE INDEX IDX1 ON ORDERS(O_TOTALPRICE);
```

Creating an index

```
CREATE INDEX IDX2 ON CUSTOMER(C_ACCTBAL);
```

SELECT statement with JOIN operation and a hint

```
SELECT /*+ INDEX_ASC(ORDERS IDX1) INDEX_ASC(CUSTOMER IDX2) */ *
FROM ORDERS JOIN CUSTOMER
      ON ORDERS.O_TOTALPRICE > CUSTOMER.C_ACCTBAL;
```

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|-----|-------------------------------------|----------|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 20G | 5015G | | 549K (10) | 00:00:22 |
| 1 | MERGE JOIN | | 20G | 5015G | | 549K (10) | 00:00:22 |
| 2 | SORT JOIN | | 450K | 46M | | 451K (1) | 00:00:18 |
| 3 | TABLE ACCESS BY INDEX ROWID BATCHED | ORDERS | 450K | 46M | | 451K (1) | 00:00:18 |
| 4 | INDEX FULL SCAN | IDX1 | 450K | | | 1069 (1) | 00:00:01 |
| * 5 | SORT JOIN | | 45000 | 6987K | 15M | 46654 (1) | 00:00:02 |
| 6 | TABLE ACCESS BY INDEX ROWID BATCHED | CUSTOMER | 45000 | 6987K | | 45079 (1) | 00:00:02 |
| 7 | INDEX FULL SCAN | IDX2 | 45000 | | | 102 (0) | 00:00:01 |

```
5 - access(INTERNAL_FUNCTION("ORDERS"."O_TOTALPRICE")>INTERNAL_FUNCTION("CUSTOMER"."C_ACCTBAL"))
      filter(INTERNAL_FUNCTION("ORDERS"."O_TOTALPRICE")>INTERNAL_FUNCTION("CUSTOMER"."C_ACCTBAL"))
```

TOP

Created by Janusz R. Getta, CSCI317 Database Performance Tuning, SIM, Session 3, 2022

14/37

SQL Tuning (2)

Outline

Index (nested loop) joins

Hash joins

Sort-merge joins

Order of join arguments in 2-way joins

Order of join arguments in m-way joins

Comparison of join algorithms

Optimization of joins through parameters

Outer joins

Nested queries

Order of join arguments in 2-way join

When using **hash join** a **smaller argument must be first**

SELECT statement with JOIN operation and a hint

```
SELECT /*+ LEADING(LINEITEM) */ *
FROM REGION JOIN LINEITEM
      ON REGION.R_COMMENT = LINEITEM.L_COMMENT;
```

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|-----|-------------------|----------|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 9 | 1998 | | 20462 (1) | 00:00:01 |
| * 1 | HASH JOIN | | 9 | 1998 | 235M | 20462 (1) | 00:00:01 |
| 2 | TABLE ACCESS FULL | LINEITEM | 1800K | 214M | | 8788 (1) | 00:00:01 |
| 3 | TABLE ACCESS FULL | REGION | 5 | 485 | | 9 (0) | 00:00:01 |

Order of join arguments in 2-way join

When using **hash join** a **smaller argument must be first**

SELECT statement with JOIN operation and a hint

```
SELECT /*+ LEADING(REGION) */ *
FROM REGION JOIN LINEITEM
      ON REGION.R_COMMENT = LINEITEM.L_COMMENT;
```

| Id | Operation | Name | Rows | Bytes | Cost | (%CPU) | Time |
|-----|-------------------|----------|-------|-------|------|--------|----------|
| 0 | SELECT STATEMENT | | 9 | 1998 | 8801 | (1) | 00:00:01 |
| * 1 | HASH JOIN | | 9 | 1998 | 8801 | (1) | 00:00:01 |
| 2 | TABLE ACCESS FULL | REGION | 5 | 485 | 9 | (0) | 00:00:01 |
| 3 | TABLE ACCESS FULL | LINEITEM | 1800K | 214M | 8788 | (1) | 00:00:01 |

1 - access("REGION"."R_COMMENT"="LINEITEM"."L_COMMENT")

SQL Tuning (2)

Outline

Index (nested loop) joins

Hash joins

Sort-merge joins

Order of join arguments in 2-way joins

Order of join arguments in m-way joins

Comparison of join algorithms

Optimization of joins through parameters

Outer joins

Nested queries

Order of join arguments in m-way join

When joining 3 or more tables join smaller tables first

SELECT statement with JOIN operation and a hint

```
SELECT /*+ LEADING(LINEITEM) */ *
FROM LINEITEM JOIN ORDERS
      ON LINEITEM.L_COMMENT = ORDERS.O_COMMENT
JOIN REGION
      ON ORDERS.O_COMMENT = REGION.R_COMMENT;
```

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|-----|-------------------|----------|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 9 | 2979 | | 24989 (1) | 00:00:01 |
| * 1 | HASH JOIN | | 9 | 2979 | | 24989 (1) | 00:00:01 |
| 2 | TABLE ACCESS FULL | REGION | 5 | 485 | | 9 (0) | 00:00:01 |
| * 3 | HASH JOIN | | 787K | 175M | 235M | 24978 (1) | 00:00:01 |
| 4 | TABLE ACCESS FULL | LINEITEM | 1800K | 214M | | 8788 (1) | 00:00:01 |
| 5 | TABLE ACCESS FULL | ORDERS | 450K | 46M | | 1950 (1) | 00:00:01 |

1 - access("ORDERS"."O_COMMENT"="REGION"."R_COMMENT")

3 - access("LINEITEM"."L_COMMENT"="ORDERS"."O_COMMENT")

Order of join arguments in m-way join

When joining **3 or more tables** join **smaller tables first**

SELECT statement with JOIN operation and a hint

```
SELECT /*+ LEADING(ORDERS) */ *
FROM LINEITEM JOIN ORDERS
      ON LINEITEM.L_COMMENT = ORDERS.O_COMMENT
JOIN REGION
      ON ORDERS.O_COMMENT = REGION.R_COMMENT;
```

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|-----|-------------------|---------|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 9 | 2979 | | 13327 (1) | 00:00:01 |
| * 1 | HASH JOIN | | 9 | 2979 | | 13327 (1) | 00:00:01 |
| * 2 | HASH JOIN | | 5 | 1030 | 51M | 4535 (1) | 00:00:01 |
| 3 | TABLE ACCESS FULL | ORDERS | 450K | 46M | | 1950 (1) | 00:00:01 |
| 4 | TABLE ACCESS FULL | REGION | 5 | 485 | | 9 (0) | 00:00:01 |
| 5 | TABLE ACCESS FULL | LINEITE | 1800K | 214M | | 8788 (1) | 00:00:01 |

1 - access("LINEITEM"."L_COMMENT"="ORDERS"."O_COMMENT")

2 - access("ORDERS"."O_COMMENT"="REGION"."R_COMMENT")

Order of join arguments in m-way join

When **joining 3 or more tables** join **smaller tables first**

SELECT statement with JOIN operation and a hint

```
SELECT /*+ LEADING(REGION) */ *
FROM LINEITEM JOIN ORDERS
      ON LINEITEM.L_COMMENT = ORDERS.O_COMMENT
JOIN REGION
      ON ORDERS.O_COMMENT = REGION.R_COMMENT;
```

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
|-----|-------------------|----------|-------|-------|-------------|----------|
| 0 | SELECT STATEMENT | | 9 | 2979 | 10752 (1) | 00:00:01 |
| * 1 | HASH JOIN | | 9 | 2979 | 10752 (1) | 00:00:01 |
| * 2 | HASH JOIN | | 5 | 1030 | 1960 (1) | 00:00:01 |
| 3 | TABLE ACCESS FULL | REGION | 5 | 485 | 9 (0) | 00:00:01 |
| 4 | TABLE ACCESS FULL | ORDERS | 450K | 46M | 1950 (1) | 00:00:01 |
| 5 | TABLE ACCESS FULL | LINEITEM | 1800K | 214M | 8788 (1) | 00:00:01 |

1 - access("LINEITEM"."L_COMMENT"="ORDERS"."O_COMMENT")

2 - access("ORDERS"."O_COMMENT"="REGION"."R_COMMENT")

SQL Tuning (2)

Outline

Index (nested loop) joins

Hash joins

Sort-merge joins

Order of join arguments in 2-way joins

Order of join arguments in m-way joins

Comparison of join algorithms

Optimization of joins through parameters

Outer joins

Nested queries

Comparison of join algorithms

Sort-merge/hash versus index-based join

- Index-based join provides a better response time
- Sort-merge/hash join provides better throughput
- Sort-merge/hash join is faster for large tables
- Sort-merge/hash join needs more cpu and memory
- Index-based join needs an index
- Hash join is superior to sort-merge join
- Index-based join is superior to sort-merge/hash join when a small number of rows is returned

SQL Tuning (2)

Outline

Index (nested loop) joins

Hash joins

Sort-merge joins

Order of join arguments in 2-way joins

Order of join arguments in m-way joins

Comparison of join algorithms

Optimization of joins through parameters

Outer joins

Nested queries

Optimization of hash join through parameters

`HASH_AREA_SIZE` system initialization parameter determines the amount of memory available to `hash join` for the creation and storage of hash table

`HASH_MULTIBLOCK_IO_COUNT` system initialization parameter determines the total number of blocks that will be written to, or read from `hash join` partition in a single I/O operation

SQL Tuning (2)

Outline

Index (nested loop) joins

Hash joins

Sort-merge joins

Order of join arguments in 2-way joins

Order of join arguments in m-way joins

Comparison of join algorithms

Optimization of joins through parameters

Outer joins

Nested queries

Outer joins

Performance of **outer join** is comparable to performance of **equijoin**

Outer join imposes a particular order of the arguments: a table which returns all rows must be used in outer loop

SQL Tuning (2)

Outline

Index (nested loop) joins

Hash joins

Sort-merge joins

Order of join arguments in 2-way joins

Order of join arguments in m-way joins

Comparison of join algorithms

Optimization of joins through parameters

Outer joins

Nested queries

Nested queries

A subquery that returns only one value is computed first

Nested SELECT statement with a subquery that returns one value

```
SELECT COUNT(*)
FROM ORDERS
WHERE ORDERS.O_TOTALPRICE = ( SELECT MIN(C_ACCTBAL)
                               FROM CUSTOMER );
```

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
|-----|-------------------|----------|-------|-------|-------------|----------|
| 0 | SELECT STATEMENT | | 1 | 6 | 2232 (1) | 00:00:01 |
| 1 | SORT AGGREGATE | | 1 | 6 | | |
| * 2 | TABLE ACCESS FULL | ORDERS | 1 | 6 | 1949 (1) | 00:00:01 |
| 3 | SORT AGGREGATE | | 1 | 6 | | |
| 4 | TABLE ACCESS FULL | CUSTOMER | 45000 | 263K | 282 (0) | 00:00:01 |

```
2 - filter("ORDERS"."O_TOTALPRICE"= (SELECT MIN("C_ACCTBAL") FROM
    "CUSTOMER" "CUSTOMER"))
```

Nested queries

If a subquery returns more than one value then a **semijoin** is used

Nested SELECT statement with a subquery that returns more than one row

```
SELECT COUNT(*)
FROM ORDERS
WHERE ORDERS.O_TOTALPRICE IN ( SELECT C_ACCTBAL
                                FROM CUSTOMER );
```

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
|-----|----------------------|----------|-------|-------|-------------|----------|
| 0 | SELECT STATEMENT | | 1 | 12 | 2233 (1) | 00:00:01 |
| 1 | SORT AGGREGATE | | 1 | 12 | | |
| * 2 | HASH JOIN RIGHT SEMI | | 40406 | 473K | 2233 (1) | 00:00:01 |
| * 3 | TABLE ACCESS FULL | CUSTOMER | 40910 | 239K | 283 (1) | 00:00:01 |
| 4 | TABLE ACCESS FULL | ORDERS | 450K | 2636K | 1949 (1) | 00:00:01 |

2 – access("ORDERS"."O_TOTALPRICE"="C_ACCTBAL")

3 – filter("C_ACCTBAL">=0)

Nested queries

If a subquery returns more than one value then a **antijoin** is used

Nested SELECT statement with NOT IN relation and a subquery that returns more than one row

```
SELECT COUNT(*)
FROM ORDERS
WHERE ORDERS.O_TOTALPRICE NOT IN ( SELECT C_ACCTBAL
                                   FROM CUSTOMER );
```

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
|-----|----------------------|----------|-------|-------|-------------|----------|
| 0 | SELECT STATEMENT | | 1 | 12 | 3087 (1) | 00:00:01 |
| 1 | SORT AGGREGATE | | 1 | 12 | | |
| * 2 | HASH JOIN RIGHT ANTI | | 441K | 5177K | 3087 (1) | 00:00:01 |
| * 3 | TABLE ACCESS FULL | CUSTOMER | 40910 | 239K | 390 (1) | 00:00:01 |
| 4 | TABLE ACCESS FULL | ORDERS | 450K | 2636K | 2696 (1) | 00:00:01 |

2 – access("ORDERS"."O_TOTALPRICE"="C_ACCTBAL")

3 – filter("C_ACCTBAL">=0)

Nested queries

If an outer and inner query operates on the same table then it is possible to optimize entire query with a single index

CREATE INDEX statement

```
CREATE INDEX IDX ON ORDERS(O_TOTALPRICE);
```

Nested SELECT statement with a subquery that returns one row

```
SELECT COUNT(*)
FROM ORDERS
WHERE ORDERS.O_TOTALPRICE = ( SELECT MIN(ORDERS.O_TOTALPRICE)
                              FROM ORDERS );
```

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
|-----|---------------------------|------|------|-------|-------------|----------|
| 0 | SELECT STATEMENT | | 1 | 6 | 3 (0) | 00:00:01 |
| 1 | SORT AGGREGATE | | 1 | 6 | | |
| * 2 | INDEX RANGE SCAN | IDX | 1 | 6 | 3 (0) | 00:00:01 |
| 3 | SORT AGGREGATE | | 1 | 6 | | |
| 4 | INDEX FULL SCAN (MIN/MAX) | IDX | 1 | 6 | 3 (0) | 00:00:01 |

2 – access("ORDERS"."O_TOTALPRICE"= (SELECT MIN("ORDERS"."O_TOTALPRICE")
FROM "ORDERS" "ORDERS"))

3 – filter("C_ACCTBAL">=0)

Nested queries

Nested query with **EXISTS** quantifier is implemented with a **semijoin** operation

SELECT statement with EXISTS quantifier

```
SELECT *
FROM CUSTOMER
WHERE EXISTS ( SELECT *
                FROM ORDERS
                WHERE CUSTOMER.C_CUSTKEY = ORDERS.O_CUSTKEY );
```

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|-----|----------------------|----------|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 30162 | 4830K | | 2960 (1) | 00:00:01 |
| * 1 | HASH JOIN RIGHT SEMI | | 30162 | 4830K | 7472K | 2960 (1) | 00:00:01 |
| * 2 | TABLE ACCESS FULL | ORDERS | 450K | 2197K | | 1949 (1) | 00:00:01 |
| 3 | TABLE ACCESS FULL | CUSTOMER | 45000 | 6987K | | 283 (1) | 00:00:01 |

- 1 - access("CUSTOMER"."C_CUSTKEY"="ORDERS"."O_CUSTKEY")
- 2 - filter("ORDERS"."O_CUSTKEY">=0)

Nested queries

Performance of a nested query with **EXISTS** quantifier implemented with **semijoin** can be improved with an **index**

CREATE INDEX statement

```
CREATE INDEX IDX ON ORDERS(O_CUSTKEY);
```

SELECT statement with EXISTS quantifier

```
SELECT *
FROM CUSTOMER
WHERE EXISTS ( SELECT *
                FROM ORDERS
                WHERE CUSTOMER.C_CUSTKEY = ORDERS.O_CUSTKEY );
```

| | Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|---|----|----------------------|----------|-------|-------|---------|-------------|----------|
| | 0 | SELECT STATEMENT | | 30162 | 4830K | | 1283 (1) | 00:00:01 |
| * | 1 | HASH JOIN RIGHT SEMI | | 30162 | 4830K | 7472K | 1283 (1) | 00:00:01 |
| * | 2 | INDEX FAST FULL SCAN | IDX | 450K | 2197K | | 272 (1) | 00:00:01 |
| | 3 | TABLE ACCESS FULL | CUSTOMER | 45000 | 6987K | | 283 (1) | 00:00:01 |

1 - access("CUSTOMER"."C_CUSTKEY"="ORDERS"."O_CUSTKEY")

2 - filter("ORDERS"."O_CUSTKEY">=0)

Nested queries

Nested query with **NOT EXISTS** negated quantifier is implemented with an **antijoin** operation

SELECT statement with NOT EXISTS negated quantifier

```
SELECT *
FROM CUSTOMER
WHERE NOT EXISTS ( SELECT *
                   FROM ORDERS
                   WHERE CUSTOMER.C_CUSTKEY = ORDERS.O_CUSTKEY );
```

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|-----|----------------------|----------|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 14838 | 2376K | | 2960 (1) | 00:00:01 |
| * 1 | HASH JOIN RIGHT ANTI | | 14838 | 2376K | 7472K | 2960 (1) | 00:00:01 |
| * 2 | TABLE ACCESS FULL | ORDERS | 450K | 2197K | | 1949 (1) | 00:00:01 |
| 3 | TABLE ACCESS FULL | CUSTOMER | 45000 | 6987K | | 283 (1) | 00:00:01 |

```
1 - access("CUSTOMER"."C_CUSTKEY"="ORDERS"."O_CUSTKEY")
2 - filter("ORDERS"."O_CUSTKEY">=0)
```

Nested queries

Performance of a nested query with **NOT EXISTS** negated quantifier implemented with a **antijoin** can be improved with an **index**

CREATE INDEX statement

```
CREATE INDEX IDX ON ORDERS(O_CUSTKEY);
```

SELECT statement with NOT EXISTS negated quantifier

```
SELECT *
FROM CUSTOMER
WHERE NOT EXISTS ( SELECT *
                   FROM ORDERS
                   WHERE CUSTOMER.C_CUSTKEY = ORDERS.O_CUSTKEY );
```

| | Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|---|----|----------------------|----------|-------|-------|---------|-------------|----------|
| | 0 | SELECT STATEMENT | | 14838 | 2376K | | 1283 (1) | 00:00:01 |
| * | 1 | HASH JOIN RIGHT ANTI | | 14838 | 2376K | 7472K | 1283 (1) | 00:00:01 |
| * | 2 | INDEX FAST FULL SCAN | IDX | 450K | 2197K | | 272 (1) | 00:00:01 |
| | 3 | TABLE ACCESS FULL | CUSTOMER | 45000 | 6987K | | 283 (1) | 00:00:01 |

1 - access("CUSTOMER"."C_CUSTKEY"="ORDERS"."O_CUSTKEY")

2 - filter("ORDERS"."O_CUSTKEY">=0)

References

T. Hasler, Expert Oracle SQL Optimization, Deployment, and Statistics ,
Apress, 2014

S. R. Alapati, D. Kuhn, B. Padfield, Oracle Database 12c Performance
Tuning Recipes A Problem-Solution Approach, Apress, 2013

G. Harrison, Oracle Performance Survival Guide, A Systematic Approach
to Database Optimization, Prentice Hall Professional Oracle Series, 2010