

## CSCI317 Database Performance Tuning

# Performance Tuning of Relational Database Server

Dr Janusz R. Getta

School of Computing and Information Technology -  
University of Wollongong

# Performance Tuning of Relational Database Server

## Outline

### Top 10 mistakes

Tuning data buffer cache

Automatic shared memory management

Tuning redo log buffer

Tuning library cache

Tuning dictionary cache

Tuning Process Global Area

Tuning I/O

Choosing data block size

Tuning operating system scheduling

# Top 10 mistakes

## (1) Bad management of database connections

- A database application always connects and disconnects to perform each operation on a database

## (2) Bad use of cursors and Shared Pool

- Not using cursors and bind variables increases time spent on parsing and optimization, dynamic **SQL** has a negative impact on performance

## (3) Bad **SQL**

- Processing of **SQL** statements (especially **SELECT**) takes more time than processing the same applications with algorithmic codes, relational tables are scanned too many times

## (4) Use of nonstandard initialization parameters

- Setting the unusual values of initialization parameters, using undocumented initialization parameters

## (5) Getting database I/O wrong

- Incorrectly distributing database files over available disk drives

# Top 10 mistakes

## (6) Redo log setup problems

- Setting too small redo log buffer, using too small number of too small redo log files

## (7) Lack of freelists, shortage of rollback segments

- Typical for **INSERT**-heavy applications, too large number of users, no enough rollback segments

## (8) Long full table scans

- Usually, caused by lack of indexing, poor transaction design, poor **SQL** optimization

## (9) High amount of "recursive" **SQL**

- Too many accesses to data dictionary caused by too many space management activities, e.g. storage allocation

## (10) Deployment and migration errors

- Usually caused by incomplete migration of relational tables from development environment or legacy systems, e.g. missing indexes, missing statistics from **ANALYZE** statement

# Performance Tuning of Relational Database Server

## Outline

Top 10 mistakes

Tuning data buffer cache

Tuning redo log buffer

Tuning library cache

Tuning dictionary cache

Tuning Process Global Area

Tuning I/O

Choosing data block size

Tuning operating system scheduling

# Tuning data buffer cache

How large data buffer cache supposed to be ?

## Consistent gets

- Number of times a block was acquired in a consistent mode, that is with the correct timestamp to provide read consistency
- It is incremented by 1 for each block read during full table scans
- It is incremented by (index height + 2\*index key entries) for indexed table lookups
- It is incremented by 1 for each block read during and index-only lookup

## Db block gets

- Number of blocks read for update
- It includes updates to blocks in **UNDO** and temporary tablespaces
- It is also incremented during extent allocation and when update to high water mark takes place

# Tuning data buffer cache

## Logical Reads (LR#)

- $LR\# = \text{consistent gets} + \text{db block gets}$

## Physical Reads (PR#)


- Number of requests to an operating system to read a database block into data buffer cache
- If a parameter `DB_FILE_MULTIBLOCK_READ_COUNT` is set to more than 1 then the **physical reads** number is incremented only by one
- Reading from temporary tablespace does not increment the number

## Data block buffer hit ratio (H)

- $H = ((LR\# - PR\#) / LR\#) * 100\%$
- In a well tuned **OLTP** system  $H \geq 95\%$  and  $H \leq 85\%$  in **OLAP** system

# Tuning data buffer cache

Physical I/O versus data buffer cache size





# Tuning data buffer cache

What should be kept where ?

**KEEP** pool

- **KEEP** pool should contain the tables that are frequently referenced by the applications, like small tables with frequent full scans, lookup tables, and data objects that are normally more than 80% cached
- The contents of **KEEP** pool are not automatically removed or loaded
- **KEEP** pool should always have data block buffer ratio equal to 100%
- The size of **KEEP** pool should be greater by 20% than the total size of data objects kept there

# Tuning data buffer cache

What should be kept where ?

## RECYCLE pool

- **RECYCLE** pool should contain the large relational tables which have rare full-table scans such that their data blocks are unlikely to be re-read
- Storing rarely and fully scanned relational tables in **RECYCLE** pool prevents more frequently used tables and indexes from being flushed out of the system
- The contents of **RECYCLE** pool are immediately removed after being used

# Tuning data buffer cache

What should be kept where ?

DEFAULT pool

- **DEFAULT** pool contains all other data objects that are not assigned to either **KEEP** or **RECYCLE** pools

# Tuning data buffer cache

What should be kept where ?

**nK** pools

- The significant performance gains can be achieved by switching to a longer data block size
- the tables with long rows have less chained rows
- full table scans and tables with large objects (**BLOB**, **CLOB**, ...) benefit from longer blocks because of less read block operations
- index range scans gather index node sequentially, temporary tablespaces used for sorting also need less I/O operations
- On the other hand the tables with smaller rows accessed randomly should be placed in smaller data blocks because they take too much space in data buffer cache
- B-tree indexes should use the largest supported size of data blocks (**32K**)
- A table should be stored in the data blocks longer than the average row size in the table to prevent row chaining
- **TEMP** tablespace benefits from the large blocks that allow to minimize disk I/O when sorting

TOP

Created by Janusz R. Getta, CSCI317 Database Performance Tuning, SIM, Session 3, 2022

12/33

# Performance Tuning of Relational Database Server

## Outline

Top 10 mistakes

Tuning data buffer cache

Tuning redo log buffer

Tuning library cache

Tuning dictionary cache

Tuning Process Global Area

Tuning I/O

Choosing data block size

Tuning operating system scheduling

# Tuning redo log buffer

## Redo log space requests (RLS#)

- The total number of times a user process waits for space in a redo log buffer so that a redo entry can be written there

## Redo entries (RE#)

- The total number of entries written to a redo log buffer due to modifications of data blocks

## Redo log hit ratio (RLH)

- $RLH = RLS\#/RE\#$
- It is recommended to have  $RLH \leq 1/5000$
- In the large OnLine Transaction Processing (OLTP) systems the size of redo log buffer should be from 2M to 4M

# Tuning redo log buffer

## Redo log space wait time

- The total elapsed time of waiting for redo log space request; should be close to zero !

## When the system waits for redo log space ?

- The system waits for redo log space when the log writer process (**LGWR**) is waiting for a disk space
- Disk space is available when a **log file switch** occurs
- **log file switch** forces a **checkpoint**
- **Check point** means that all "dirty blocks" in data buffer cache must be written to persistent storage (disk drive)

## Consequence of too small redo log buffer

- Smaller redo log buffer forces the log writer process (**LGWR**) to write to a log file more frequently

# Tuning redo log buffer

How to detect log buffer contention ?

- A simple way to detect redo buffer contention is to analyze the latch usage statistics and in particular find a ratio of **gets** to **misses** or a ratio of **immediate\_gets** to **immediate\_misses**
- the ration should not exceed 1%



# Performance Tuning of Relational Database Server

## Outline

Top 10 mistakes

Tuning data buffer cache

Tuning redo log buffer

Tuning library cache

Tuning dictionary cache

Tuning Process Global Area

Tuning I/O

Choosing data block size

Tuning operating system scheduling

# Tuning library cache

## Pin (P#)

- **Pin** indicates that parse tree of SQL statement to be processed is included in library cache

## Reload (R#)

- **Reload** indicates a cache miss; a parse tree is not in library cache or it has been removed from a library cache

## Library cache hit ratio (LCH)

- $LCH = ( P\# / (P\# + R\#) ) * 100\%$
- It is recommended to have  $LCH \geq 99\%$

# Tuning library cache

## Principles

- Use as much generic code as possible
- Use bind variables (:v) instead of constants in SQL statements

```
SELECT *  
FROM EMPLOYEE  
WHERE E# = 25;
```

sql

```
VARIABLE EMPNUM NUMBER  
BEGIN  
  :EMPNUM := 25;  
END;  
SELECT *  
FROM EMPLOYEE  
WHERE E# = :EMPNUM;
```

sql

- Increase size of library cache when necessary
- Prevent invalidations (a data object used in SQL statement is modified with **ALTER**, **DROP**, or **ANALYZE** statement)

# Performance Tuning of Relational Database Server

## Outline

Top 10 mistakes

Tuning data buffer cache

Tuning redo log buffer

Tuning library cache

Tuning dictionary cache

Tuning Process Global Area

Tuning I/O

Choosing data block size

Tuning operating system scheduling

# Tuning dictionary cache

## Get (G#)

- **Get** indicates that relevant data from data dictionary are found in dictionary cache

## Getmiss (M#)

- **Getmiss** indicates that relevant data from data dictionary are not found in dictionary cache

## Dictionary cache hit ratio (DCH)

- $DCH = ( G\# / (G\# + M\#) ) * 100\%$
- It is recommended to have  $DCH \geq 90\%$

# Performance Tuning of Relational Database Server

## Outline

Top 10 mistakes

Tuning data buffer cache

Tuning redo log buffer

Tuning library cache

Tuning dictionary cache

Tuning Process Global Area

Tuning I/O

Choosing data block size

Tuning operating system scheduling

# Tuning Process Global Area

**PGA** memory management is controlled by system initialization parameter **WORKAREA\_SIZE\_POLICY**

In **manual PGA memory management** mode the values of system initialization parameters **SORT\_AREA\_SIZE**, **SORT\_AREA\_RETAINED\_SIZE**, and **HASH\_AREA\_SIZE** have the largest impact on **PGA** size

In **manual PGA memory management** mode the optimal size of **PGA** is equal to

$((1\text{M (Unix) or } 2\text{M Win}) + \text{sort area size} + \text{hash area size}) * \text{number of connected users through dedicated connections} + 2\text{M}$

In **automatic PGA memory management** **PGA** size is determined by system initialization parameter **PGA\_AGGREGATE\_TARGET**

It controls how much memory a server can allocated for all work areas used to sort and hash data

# Tuning Process Global Area

Then, **PGA** size = total transient memory available to the system - the maximum size of **System Global Area (SGA)**

In multithreaded configuration **PGA** consumes less memory even it consumes more shared memory as sort and hash areas are kept in shared memory)

**Automatic PGA memory management** is recommended in a system where a large number of users needs medium amounts of memory for sorting and hashing

**Manual PGA memory management** is recommended in a system where a small number of users needs large amounts of memory for sorting and hashing

**Automatic PGA memory management** is recommended for end-user sessions while **manual PGA memory management** is recommended for large batch jobs running when there is no other activities



# Performance Tuning of Relational Database Server

## Outline

Top 10 mistakes

Tuning data buffer cache

Tuning redo log buffer

Tuning library cache

Tuning dictionary cache

Tuning Process Global Area

Tuning I/O

Choosing data block size

Tuning operating system scheduling

# Tuning I/O

Tuning I/O is effective when an application is I/O-bound, for example it spends majority of time on waiting until I/O operations are completed

Tuning I/O must take under the consideration physical parameters of persistent storage devices, number of persistent storage devices, and controllers

Tuning I/O includes application of software or hardware disk stripping feature of operating system

Parameters to be determined: stripe depth (size of a single stripe, 256K .. 1M) and stripe width (number of disk drives)

Important system initialization parameters that have impact on I/O tuning

DB_BLOCK_SIZE	Determines size of database blocks
DB_FILE_MULTIBLOCK_READ_COUNT	Determines the maximum I/O size
HASH_AREA_SIZE	Determines I/O size for hash operations
SORT_AREA_SIZE	Determines I/O size for sort operations
OS block size	Determines an operating system block size

# Tuning I/O


## Tuning I/O over tables indexes and temporary tablespaces

- Spread the database files across multiple persistent storage devices
- Separate the files with a high I/O rates from the remaining files
- If the files with high I/O activities contain tables and indexes then try to tune SQL or application code
- If the files with high I/O activities implement a temporary tablespace (**TEMP**) then tune the applications performing sorting

## Tuning I/O over redo log and archived redo log files

- Place all redo log files on one disk without any other database files
- Members of the same redo log group should be located on different persistent storage devices without any other database files
- Perform stripping of redo log files across many persistent storage devices
- If a persistent storage device contains archived redo logs then no other process should compete with archiver process for the access to the device

# Tuning I/O



## Sizing redo log files

- Larger redo log size provide better performance
- Smaller redo log size increases checkpoint activity and decreases performance
- Checkpoint frequency depends on redo log file size and on parameter **FAST\_START\_MTTR\_TARGET** which restricts instance recovery time
- Oracle tries to automatically perform checkpoint to limit amount of recovery time determined by the parameter
- Optimal size of redo log files is suggested in **OPTIMAL\_LOGFILE\_SIZE** column of **V\$INSTANCE\_RECOVERY** view

# Performance Tuning of Relational Database Server

## Outline

Top 10 mistakes

Tuning data buffer cache

Tuning redo log buffer

Tuning library cache

Tuning dictionary cache

Tuning Process Global Area

Tuning I/O

Choosing data block size

Tuning operating system scheduling

# Choosing data block size

To minimize a number of **read operations**

- A block size of **8K** is considered to be the most universal solution
- If rows are small and access is random then choose a smaller block size
- If rows are small and access is sequential then choose a larger block size
- If rows are small and access is both random and sequential then choose a larger block size
- If rows are large and contain **LOBs** then choose a larger block size

To minimize a number of **write operations**

- A block size of **8K** is considered to be the most universal solution for OnLine Transaction Processing (**OLTP**) systems
- OnLine Analytical processing (**OLAP**) systems (Data Warehouses) benefit from large block size

# Performance Tuning of Relational Database Server

## Outline

Top 10 mistakes

Tuning data buffer cache

Tuning redo log buffer

Tuning library cache

Tuning dictionary cache

Tuning Process Global Area

Tuning I/O

Choosing data block size

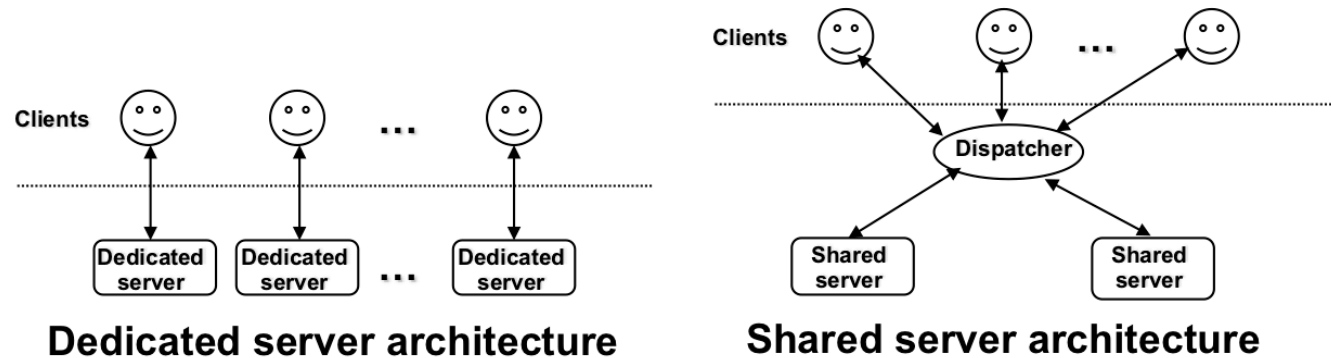
Tuning operating system scheduling

# Tuning operating system scheduling

## Principles

Minimize amount of time spent on **switching context** (activating of a different thread of control)

Choose operating system that has a lightweight thread switching facility



Minimize a number of times **context switching** occurs (minimize I/O requests, time-slice driven interrupts)



# References

Oracle 19c Database Performance Tuning Guide, 2019

T. Kyte, Expert Oracle Database Architecture, 9i, 10g, and 11g, APress, 2nd ed. 2011

G. Harrison Oracle Performance Survival Guide, Prentice Hall, 2010

D. Shasha and P. Bonnet Database Tuning Principles, Experiments, and Troubleshooting Techniques, Morgan Kaufmann, 2003, chapter 2, (UoW Library closed collection)