

CSCI317 Database Performance Tuning

In-Memory, GPU, and NVM

Database Systems

Dr Janusz R. Getta

School of Computing and Information Technology -
University of Wollongong

In-Memory, GPU, and NVM Database Systems

Outline

In-Memory Database System ? What it it ?

In-Memory Database Oracle 19c

GPU Database systems

Non-Volatile Memory Database System ? What is it ?

Non-Volatile Memory

Architecture of NVM Database Systems

In-Memory Database System ? What it it ?

The cost of transient memory has been continuously decreasing and the amount of transient memory that can be used by a database server has been continuously increasing

Many databases of modest size can now be comfortably be stored within the transient memory of a single server

An **In-Memory database (IMDB)** stores data in a computer's transient memory instead of persistent storage to produce quicker response times

A central problem is how to make data stored for a long time in transient memory safe

In a traditional database system, **COMMIT** operation the writes the modifications performed on a database to a transaction log in persistent storage

Taking full advantage of a large transient memory system requires an architecture of database system that makes data persistent despite a power failure

In-Memory Database System ? What it it ?

In-Memory database systems must address the following architectural problems

- **Cache-less architecture** because there is no point caching in memory what is already stored in memory
- **Alternative persistence model** to ensure that no data loss occurs due to power failure

In-memory databases use the following techniques to ensure that data kept in transient memory is not lost

- Replicating data on other systems
- Writing from time to time complete database images to persistent storage
- Writing out the outcomes of transactions to an append-only persistent storage

In-Memory, GPU, and NVM Database Systems

Outline

In-Memory database system ? What is it ?

In-Memory Database Oracle 19c

GPU Database systems

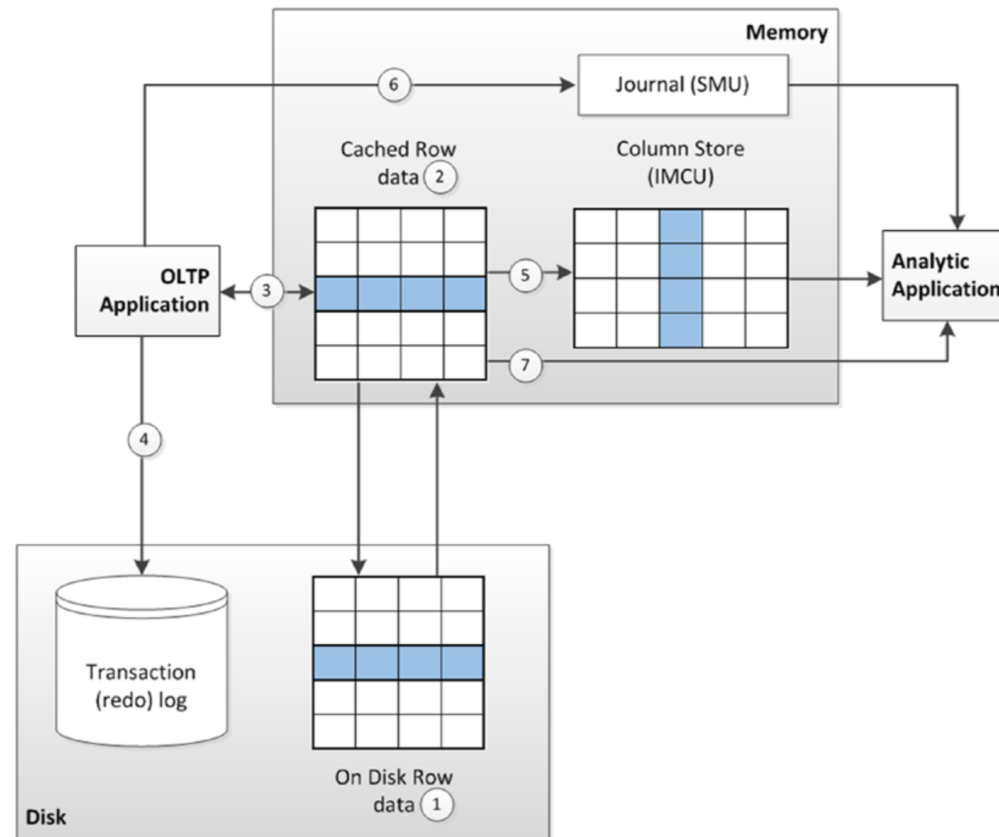
Non-Volatile Memory Database System ? What is it ?

Non-Volatile Memory

Architecture of NVM Database Systems

In-Memory Database Oracle 19c

Oracle 19c In-Memory architecture



In-Memory Database Oracle 19c

Persistence is achieved through the following actions

- (1) Data is maintained in disk files (1) and ...
- (2) ... it is cached in memory
- (3) An OLTP application primarily reads and writes from memory
- (4) Any committed transactions are written immediately to the transaction log on disk
- (5) When required or as configured, row data is loaded into a columnar representation for use by analytic applications
- (6) Any transactions that are committed once the data is loaded into columnar format are recorded in a journal
- (7) Analytic queries access the journal to determine if they need to read updated data from the row store or possibly rebuild the column

In-Memory Database Oracle 19c

In-Memory Database uses In-Memory Area located in System Global Area (SGA)

In-Memory Area is subdivided into 2 pools: 1MB pool and 64K pool

1MB pool stores columnar formatted data, 64K pool stores metadata

Listing In-Memory Area parameters

```
SELECT POOL, ALLOC_BYTES, USED_BYTES, POPULATE_STATUS  
FROM V$INMEMORY_AREA;
```

The size of In-Memory Area is controlled by the system initialization parameter INMEMORY_SIZE (default 0)

In-Memory Area must have a minimum size of 100MB

It is possible to increase In-Memory Area by increasing INMEMORY_SIZE parameter with ALTER SYSTEM command (assuming that there is spare memory in SGA)

Changing INMEMORY_SIZE parameter

```
ALTER SYSTEM SET INMEMORY_SIZE=300M SCOPE=SPFILE;
```


In-Memory Database Oracle 19c

Note, that the system must be shutdown and restarted when an option **SCOPE=SPFILE** is used !

| Values of In-Memory Area parameters | | |
|-------------------------------------|-------------|------------|
| P00L | ALLOC_BYTES | USED_BYTES |
| ----- | ----- | ----- |
| 1MB P00L | 217055232 | 0 |
| 64KB P00L | 79691776 | 0 |

INMEMORY attribute can be specified for tablespace, table, partition, subpartition, or materialized view

```
ALTER TABLESPACE USERS DEFAULT INMEMORY;
```

Enabling INMEMORY attribute for a tablespace

Populating a table **LINEITEM** in In-Memory Area

```
ALTER TABLE LINEITEM INMEMORY;
```

Enabling INMEMORY attribute for a table

In-Memory Database Oracle 19c

Explaining query processing plan of **SELECT** statement that uses **LINEITEM** table

Query processing plan for In-Memory table

```
EXPLAIN PLAN FOR SELECT L_EXTENDEDPRICE FROM LINEITEM;
```

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
|----|----------------------------|----------|-------|-------|-------------|----------|
| 0 | SELECT STATEMENT | | 1800K | 10M | 12148 (1) | 00:00:01 |
| 1 | TABLE ACCESS INMEMORY FULL | LINEITEM | 1800K | 10M | 12148 (1) | 00:00:01 |

Excluding a column **O_TOTALPRICE** in a table **ORDERS** from In-Memory Area

Excluding a column from In-Memory Area

```
ALTER TABLE ORDERS INMEMORY NO INMEMORY(O_TOTALPRICE);
```

Excluding a table **LINEITEM** from In-Memory Area

Excluding a table from In-Memory Area

```
ALTER TABLE LINEITEM NO INMEMORY;
```

In-Memory, GPU, and NVM Database Systems

Outline

In-Memory database system ? What is it ?

In-Memory Database Oracle 19c

GPU Database systems

Non-Volatile Memory Database System ? What is it ?

Non-Volatile Memory

Architecture of NVM Database Systems

GPU Database systems

Graphical Processing Unit (GPU) database systems use GPU chipset technology for parallel processing of data analytics queries mainly computed at Data Warehousing systems

GPU database systems have been on the market for a few years and slowly became adopted in the high-performance computing areas

Most GPU databases have a master-slave architecture where a central CPU-based node submits the subqueries in parallel to an array of GPU-accelerated database instances, each on a separate server

The individual servers execute their subqueries in parallel, send and result sets back to the master node

The master node then combines them and sends the final result back to the client

GPU database system works best when you have data analytics applications that can be easily broken down into parallelized tasks

GPU Database systems

The existing systems

- **Brytlyt**: built on PostgreSQL, can cost-effectively query multibillion row datasets in seconds, easily scalable, supports additions and removals of GPU-acceleration nodes, optimized for real-time insights on large and streaming data sets
- **BlazingDB**: processes simple SQL queries on massive data sets, optimized for data warehousing workloads, its SQL engine runs on a cluster of distributed GPU servers
- **Kinetica**: uses in-memory storage and distributed processing, up to 6,000 cores in parallel, processes SQL queries on billions of rows in microseconds, provides native support for geospatial objects and it comes with a suite of geospatial functions
- **MapD**: uses in-memory storage, compiles SQL queries with a just-in-time **LLVM**-based compiler into machine code that can run on Nvidia GPUs as well as X86 or Power CPUs, parallelizes computation across multiple GPUs and CPUs
- **SQream**: SQL based column-oriented database system, processes trillions of rows in near real-time, uses CPUs as well as massively parallel multi-core GPU nodes, uses automatic and transparent partitioning, allows for selective access to the required subset of columns

In-Memory, GPU, and NVM Database Systems

Outline

In-Memory database system ? What is it ?

In-Memory Database Oracle 19c

GPU Database systems

Non-Volatile Memory Database System ? What is it ?

Non-Volatile Memory

Architecture of NVM Database Systems

Non-Volatile Memory Database System ? What is it ?

Non-Volatile Memory (NVM) is a name for a group of technologies such as **Phase-Change RAM (PCRAM)**, **Magnetic RAM (MRAM)** and **Resistive RAM (RRAM)** that enable memory chips that are non-volatile (persistent), require low energy, and have density and latency closer to current DRAM chips

A commercial implementation of **NVM** is the **Intel® Optane™ DC Persistent Memory Module** (simply called as **Optane Memory**, March 2019)

NVM has 4 times faster **Input/Output Operations Per Second (IOPs)** than SSD and **seek time** for data is ten times faster

NVM supports **byte-addressable** loads and stores with low latency and it can be used for efficient architecture of **In-Memory database systems**

NVM Database System is an **In-Memory Database System** implemented in **Non-Volatile Memory**

In-Memory, GPU, and NVM Database Systems

Outline

In-Memory database system ? What is it ?

In-Memory Database Oracle 19c

GPU Database systems

Non-Volatile Memory Database System ? What is it ?

Non-Volatile Memory

Architecture of NVM Database Systems

Non-Volatile Memory

Important properties of Non-Volatile Memory

- **Byte-Addressability**: NVM supports byte-addressable loads and stores, no need to transfer data in blocks
- **High Write Throughput**: NVM delivers more than an order of magnitude higher write throughput compared to SSD
- **Read-Write Asymmetry**: in certain NVM technologies, write take longer to complete when compared to read and excessive writes to a single memory cell can destroy it

| Attribute | NVM | | | | | |
|----------------------|------------|------------|---------|-----------|-------------|------------|
| | DRAM | PCM | RRAM | MRAM | SSD | HDD |
| Read latency | 60 ns | 50 ns | 100 ns | 20 ns | 25 μ s | 10 ms |
| Write latency | 60 ns | 150 ns | 100 ns | 20 ns | 300 μ s | 10 ms |
| Sequential bandwidth | 60 GB/s | 10 GB/s | 10 GB/s | 5 GB/s | 1 GB/s | 0.1 GB/s |
| \$/GB | 10 | 1 | 1 | 1 | 0.25 | 0.02 |
| Addressability | Byte | Byte | Byte | Byte | Block | Block |
| Persistent | No | Yes | Yes | Yes | Yes | Yes |
| Endurance | $>10^{16}$ | $>10^{10}$ | 10^8 | 10^{15} | 10^5 | $>10^{16}$ |

In-Memory, GPU, and NVM Database Systems

Outline

In-Memory database system ? What is it ?

In-Memory Database Oracle 19c

GPU Database systems

Non-Volatile Memory Database System ? What is it ?

Non-Volatile Memory

Architecture of NVM Database Systems

Architecture of NVM Database Systems

NVM-based architectures have several key advantages over the current in-memory systems

- **NVM**-based architecture adopts a logging and recovery protocol that improves the availability of the DBMS by 100 times when compared with the write-ahead logging protocol
- **NVM**-based storage engines utilize the durability and byte-addressability properties of **NVM** to assure better persistent storage utilization
- **NVM**-based storage engines use the index tailored for **NVM**
- Data buffer management policy uses the direct-addressability property of **NVM** to reduce data migration
- **NVM**-based architecture improves the runtime performance, availability, operational cost, and development cost of DBMS

Architecture of NVM Database Systems

There are two possible applications of **NVM** in a database system

- (1) Database system only has **NVM** storage with no **DRAM** (**NVM**-only architecture)
- (2) **NVM** is added as another level of the storage hierarchy between **DRAM** and **SSD** (**NVM**+**DRAM** architecture)

Both approaches can be applied to **in-memory oriented** database systems and **storage (SSD/HDD) oriented** database systems

The experiments show that **in-memory oriented** systems are better suited to take advantage of **NVM** and outperform **storage oriented** systems

Both **in-memory oriented** and **storage oriented** systems are doing unnecessary work in their storage and recovery methods, since all writes to **NVM** are durable

It means that neither system is ideally suited for **NVM**

Instead, a new system is needed with lightweight recovery processes designed to utilize the non-volatile property of **NVM**

References

G. Harrison, Next Generation Databases NoSQL, NewSQL, and Big Data, Apress, 2015 Chapter 6, Column Databases

Oracle Database In-Memory with Oracle Database 19c, Technical Overview, February 2019

J. Arulraj, A. Pavlo, Non-Volatile Memory Database Management Systems, Synthesis Lectures on Data Management, Morgan & Claypool, 2019