

# CSCI317 – Database Performance Tuning

*Tutorial*

[sjapit@uow.edu.au](mailto:sjapit@uow.edu.au)

*17 August 2022*



# *Query Processing Plans*

# Interpretation of query processing plans

Consider the following conceptual schema of a two-dimensional data cube.

PLAN\_TABLE\_OUTPUT

Plan hash value: 4005457340

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		449K	30M		15866 (1)	00:00:01
* 1	HASH JOIN RIGHT ANTI		449K	30M		15866 (1)	00:00:01
* 2	TABLE ACCESS FULL	LINEITEM	1	9		12152 (1)	00:00:01
* 3	HASH JOIN		450K	26M	2776K	3712 (1)	00:00:01
4	TABLE ACCESS FULL	CUSTOMER	45000	2241K		389 (0)	00:00:01
* 5	TABLE ACCESS FULL	ORDERS	450K	4833K		2696 (1)	00:00:01

Predicate Information (identified by operation id):

- 1 - access("L\_ORDERKEY"="ORDERS"."O\_ORDERKEY")
- 2 - filter("L\_TAX">20)
- 3 - access("C\_CUSTKEY"="O\_CUSTKEY")
- 5 - filter("O\_CUSTKEY">=0)

# Interpretation of query processing plans

Find and draw a syntax tree of a query processing plan listed above and discover a respective SELECT statement that may have such query processing plan.

# Query Processing Plan Explained:

- An execution plan shows the detailed steps necessary to execute a SQL statement.
- These steps are expressed as a set of database operators that consume and produce rows.
- The order of the operators and their implementations is decided by the query optimizer using a combination of query transformations and physical optimization techniques.
- While the display is commonly shown in a tabular format, the plan is in fact tree-shaped.

# For example:

PLAN\_TABLE\_OUTPUT

Plan hash value: 4005457340

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		449K	30M		15866 (1)	00:00:01
* 1	HASH JOIN RIGHT ANTI		449K	30M		15866 (1)	00:00:01
* 2	TABLE ACCESS FULL	LINEITEM	1	9		12152 (1)	00:00:01
* 3	HASH JOIN		450K	26M	2776K	3712 (1)	00:00:01
4	TABLE ACCESS FULL	CUSTOMER	45000	2241K		389 (0)	00:00:01
* 5	TABLE ACCESS FULL	ORDERS	450K	4833K		2696 (1)	00:00:01

- The tabular representation is a top-down, left-to-right traversal of the execution tree.
- When you read a plan tree you should start from the inner most indented lines (also known as the leaf nodes of the execution tree) and work across and follows with next outer indented lines.

# For example:

PLAN\_TABLE\_OUTPUT

Plan hash value: 4005457340

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		449K	30M		15866 (1)	00:00:01
* 1	HASH JOIN RIGHT ANTI		449K	30M		15866 (1)	00:00:01
* 2	TABLE ACCESS FULL	LINEITEM	1	9		12152 (1)	00:00:01
* 3	HASH JOIN		450K	26M	2776K	3712 (1)	00:00:01
4	TABLE ACCESS FULL	CUSTOMER	45000	2241K		389 (0)	00:00:01
* 5	TABLE ACCESS FULL	ORDERS	450K	4833K		2696 (1)	00:00:01

- In the above example, begin by looking at the leaves of the tree. In this case the leaves of the tree are implemented using a full table scans of the ORDERS and the CUSTOMER tables.



# For example:

PLAN\_TABLE\_OUTPUT

Plan hash value: 4005457340

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		449K	30M		15866 (1)	00:00:01
* 1	HASH JOIN RIGHT ANTI		449K	30M		15866 (1)	00:00:01
* 2	TABLE ACCESS FULL	LINEITEM	1	9		12152 (1)	00:00:01
* 3	HASH JOIN		450K	26M	2776K	3712 (1)	00:00:01
4	TABLE ACCESS FULL	CUSTOMER	45000	2241K		389 (0)	00:00:01
* 5	TABLE ACCESS FULL	ORDERS	450K	4833K		2696 (1)	00:00:01

- The rows produced by these table scans will be consumed by the join operator. Here the join operator is a hash-join (other alternatives include nested-loop or sort-merge join).



# For example:

PLAN\_TABLE\_OUTPUT

Plan hash value: 4005457340

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		449K	30M		15866 (1)	00:00:01
* 1	HASH JOIN RIGHT ANTI		449K	30M		15866 (1)	00:00:01
* 2	TABLE ACCESS FULL	LINEITEM	1	9		12152 (1)	00:00:01
* 3	HASH JOIN		450K	26M	2776K	3712 (1)	00:00:01
4	TABLE ACCESS FULL	CUSTOMER	45000	2241K		389 (0)	00:00:01
* 5	TABLE ACCESS FULL	ORDERS	450K	4833K		2696 (1)	00:00:01

- Finally the result of the hash-join is hash-join with rows produced by a full-table scan of the table LINEITEM using right-anti-join.
- The final result is then sent to the end user.



# Query Processing Plans

*Drawing the syntax tree*

# Drawing a syntax tree

PLAN\_TABLE\_OUTPUT

Plan hash value: 4005457340

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		449K	30M		15866 (1)	00:00:01
* 1	HASH JOIN RIGHT ANTI		449K	30M		15866 (1)	00:00:01
* 2	TABLE ACCESS FULL	LINEITEM	1	9		12152 (1)	00:00:01
* 3	HASH JOIN		450K	26M	2776K	3712 (1)	00:00:01
4	TABLE ACCESS FULL	CUSTOMER	45000	2241K		389 (0)	00:00:01
* 5	TABLE ACCESS FULL	ORDERS	450K	4833K		2696 (1)	00:00:01

Predicate Information (identified by operation id):

- 1 - access("L\_ORDERKEY"="ORDERS"."O\_ORDERKEY")
- 2 - filter("L\_TAX">20)
- 3 - access("C\_CUSTKEY"="O\_CUSTKEY")
- 5 - filter("O\_CUSTKEY">=0)

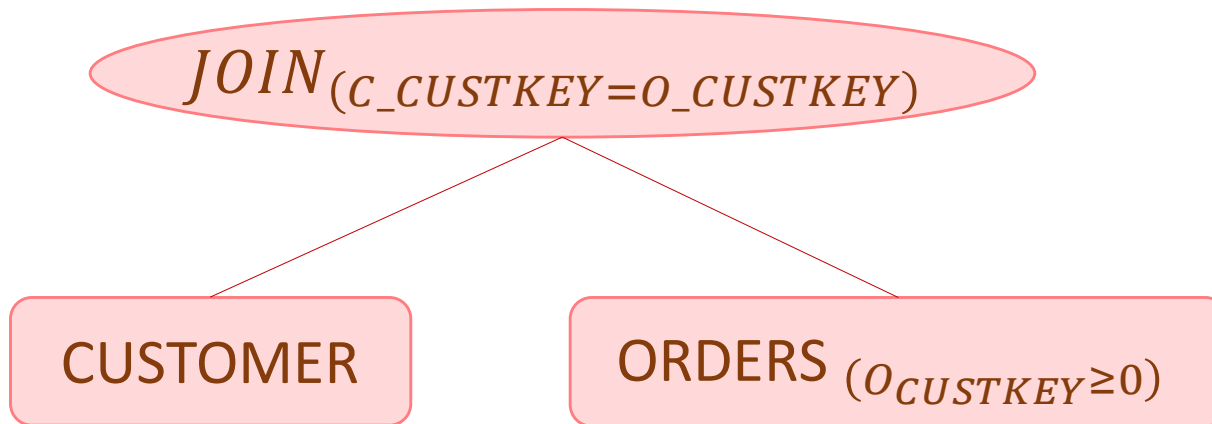
Starting from the inner most indented lines (in this example, lines 4 and 5, draw a leaf node for each line.

- Starting from the inner most indented lines (in this example, lines 4 and 5, draw a leaf node for each line.
- Indicate any filter condition if exist.

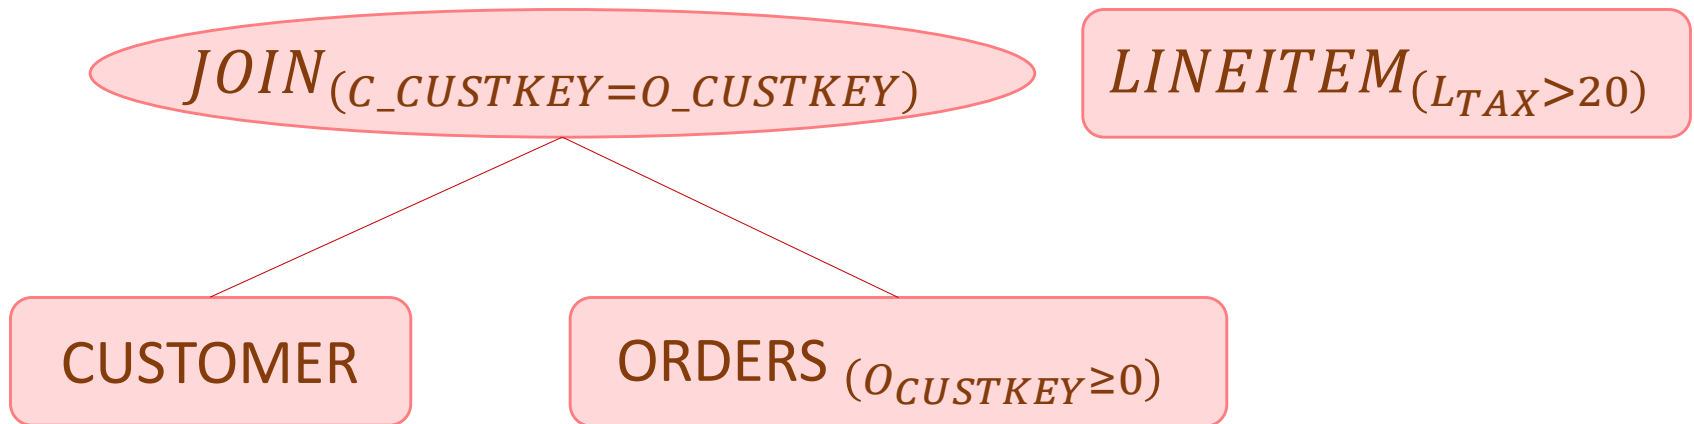
CUSTOMER

ORDERS ( $O_{CUSTKEY} \geq 0$ )

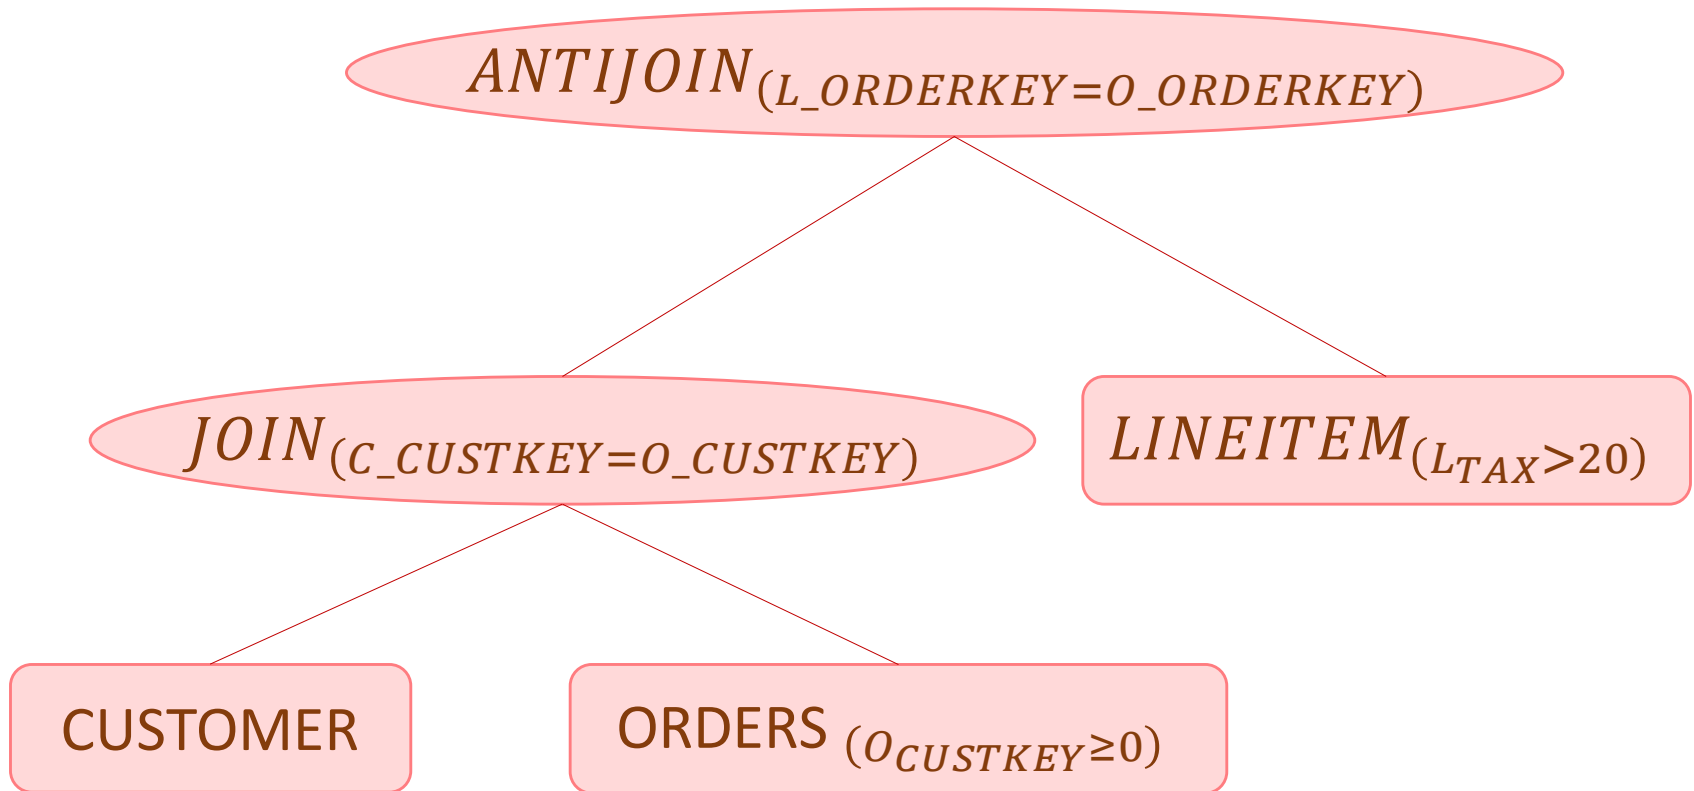
- Next draw the join operation.
- Indicate the join condition.



- Fall all operations that have the same indentation level as the join operator, draw a node for each of the operation.
- Indicate the filter condition if exist.

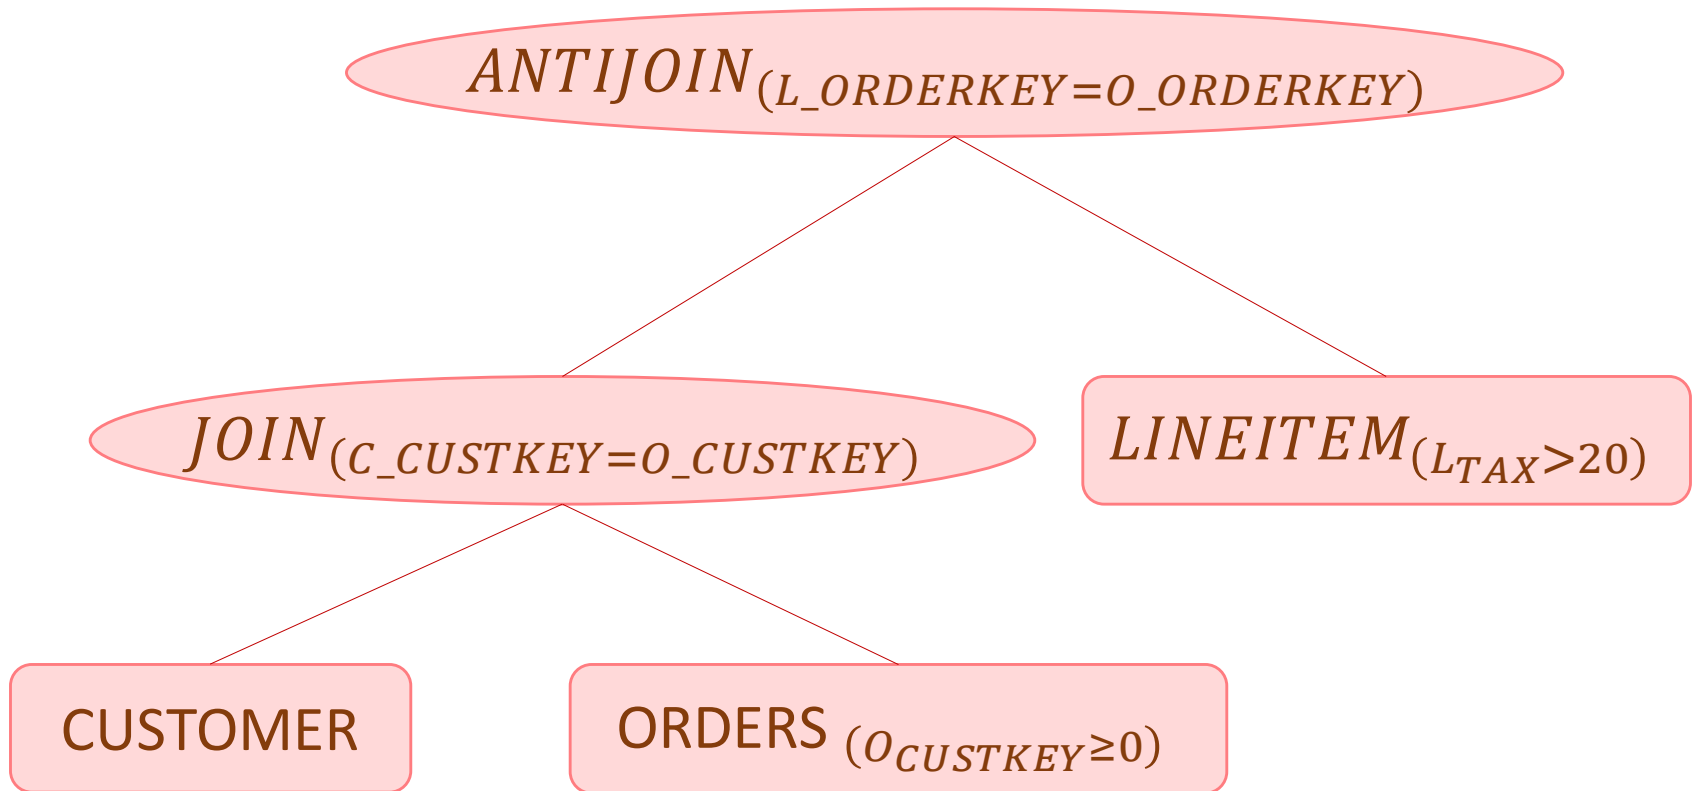


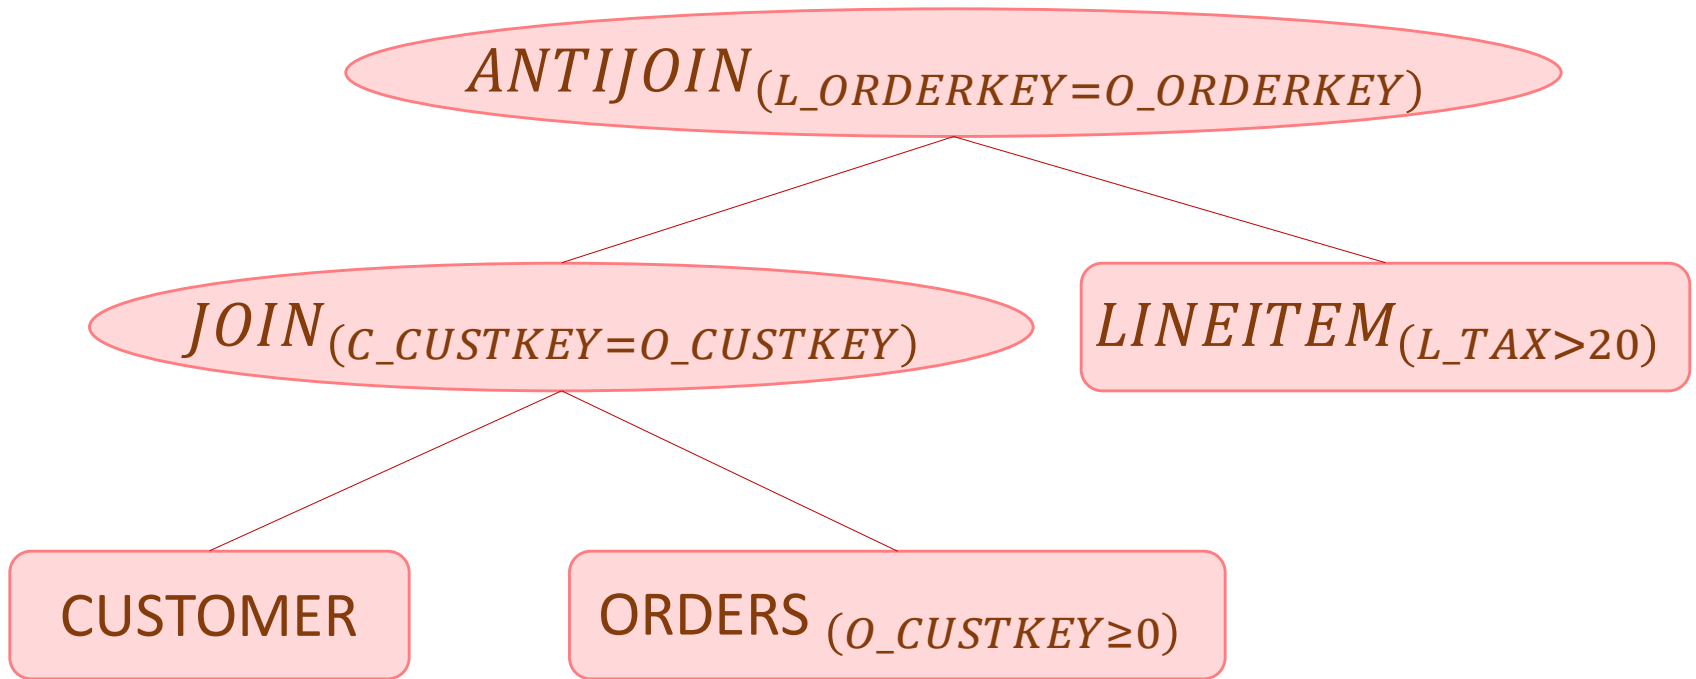
- Next draw the join operation.
- Indicate the join condition.





Stop when reach to the top (outer most indentation.)





```
SELECT *  
FROM ORDERS JOIN CUSTOMER  
    ON C_CUSTKEY = O_CUSTKEY  
WHERE O_ORDERKEY NOT IN (SELECT L_ORDERKEY  
                        FROM LINEITEM  
                        WHERE L_TAX > 20)
```